

# Music Data Analysis

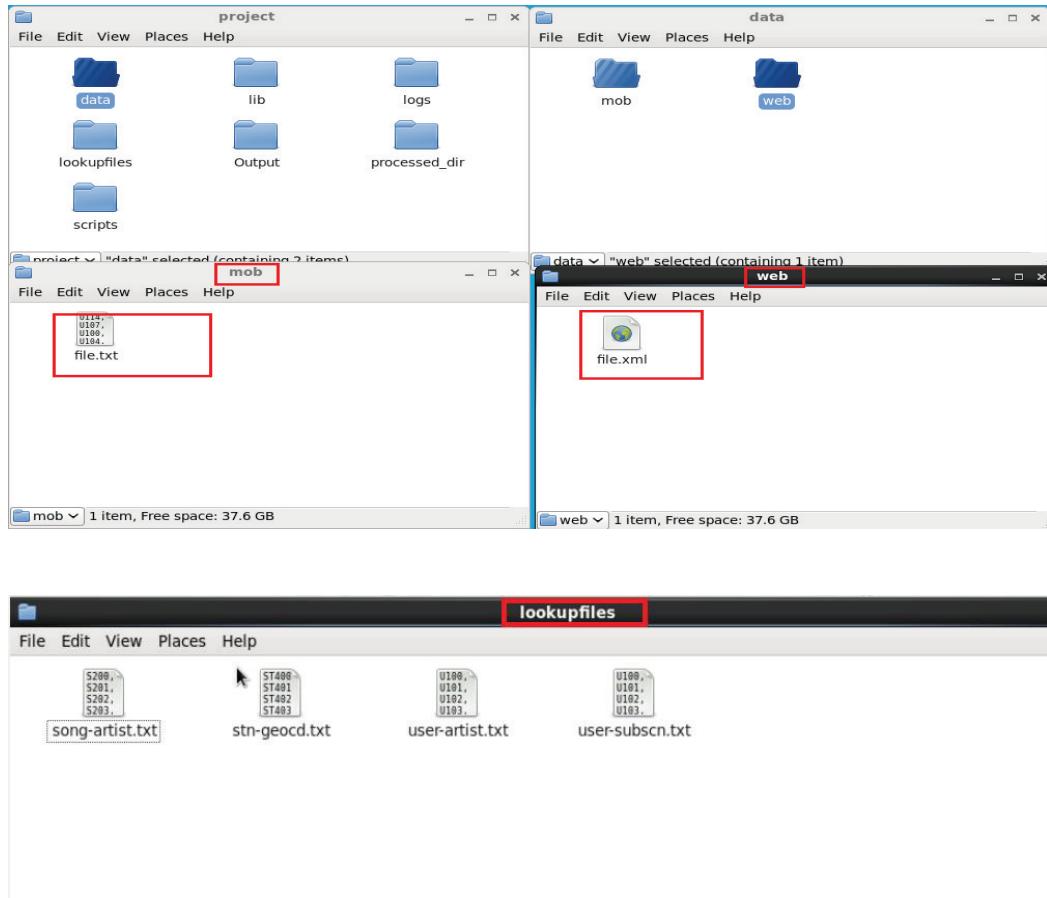
## 1. Introduction:

A leading music-catering company is planning to analyze large amount of data received from varieties of sources, namely mobile app and website to track the behaviour of users, classify users, calculate royalties associated with the song and make appropriate business strategies. The file server receives data files periodically after every 3 hours.

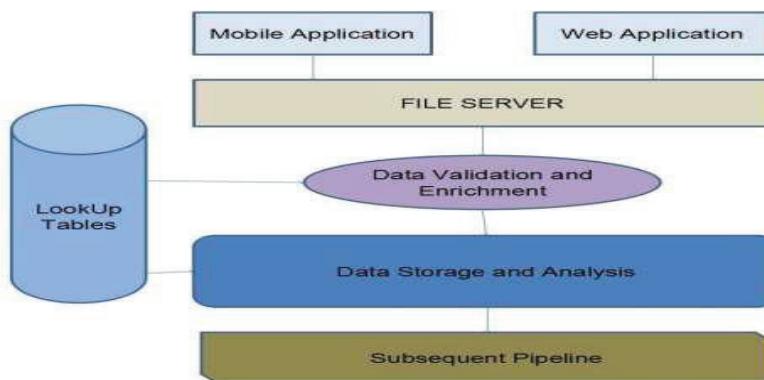
So as per periodic data for every 3 hours, we need to calculate the queries and return the results to business/customer.

## 2. DATASET:

1. Data coming from web applications reside in /data/web and has xml format.
2. Data coming from mobile applications reside in /data/mob and has csv format.
3. Data present in lookup directory should be used in HBase.

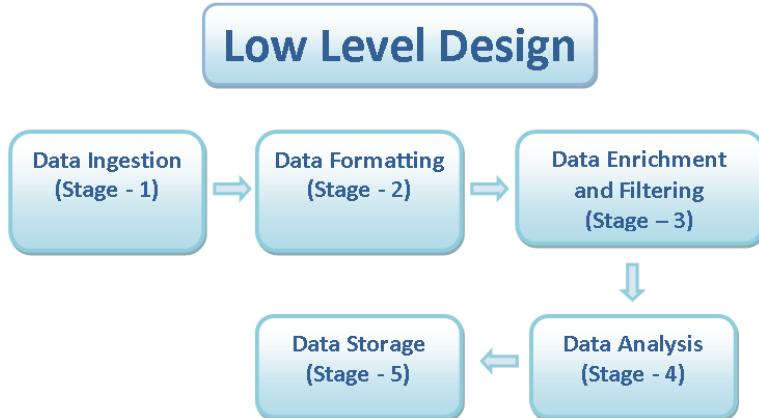


## 3. Flow of Operations:

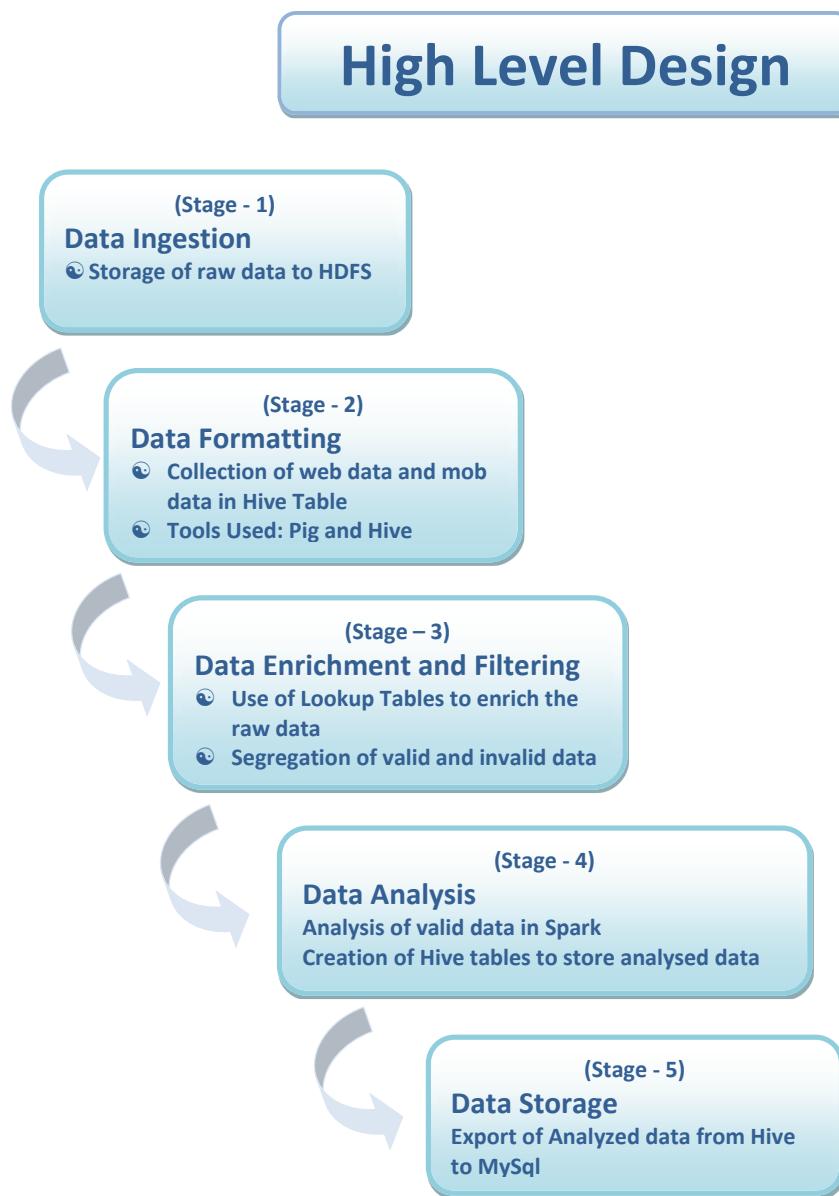


# Music Data Analysis

## 4. Low Level Design:



## 5. High Level Design



# Music Data Analysis

## 6. Implementation:

### a) Starting all daemons in Hadoop:

- In this script we will start all the daemons that are required to start the services like hive, hbase, MySQL etc.



```
#!/bin/bash
if [ -f "/home/acadgild/project/logs/current-batch.txt" ]
then
echo "Batch File Found!"
else
echo -n "1" > "/home/acadgild/project/logs/current-batch.txt"
fi
chmod 775 /home/acadgild/project/logs/current-batch.txt
batchid= cat /home/acadgild/project/logs/current-batch.txt
LOGFILE=/home/acadgild/project/logs/log_batch_$batchid
echo "Starting daemons..." |>> $LOGFILE
# To start the hadoop daemons:
/usr/local/hadoop-2.6.0/sbin/start-all.sh
# To start the Hmaster service:
/usr/local/hbase/bin/start-hbase.sh
# To start the Jobhistory server service:
mr-jobhistory-daemon.sh start historyserver
# To start mysql service:
sudo service mysqld start
# To start hive metastore:
hive --service metastore
```

- Terminal with all the daemons started screenshot.

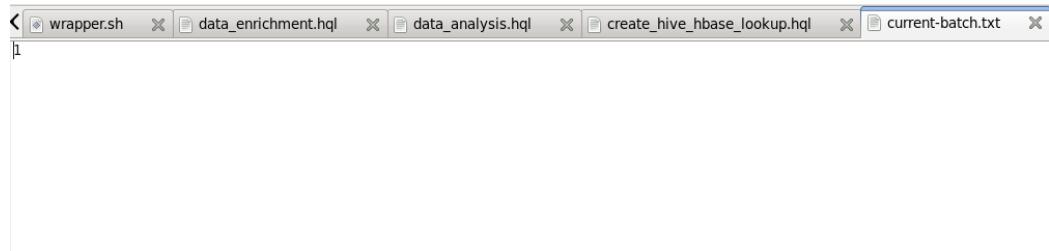
```
[acadgild@localhost scripts]$ sh start-daemons.sh
Batch file found!
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
17/09/11 12:42:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
acadgild@localhost's password:
localhost: starting namenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-namenode-localhost.localdomain.out
acadgild@localhost's password:
localhost: starting datanode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-datanode-localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
acadgild@0.0.0.0's password:
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-secondarynamenode-localhost.localdomain.out
17/09/11 12:43:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-resourcemanager-localhost.localdomain.out
acadgild@localhost's password:
localhost: starting nodemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-nodemanager-localhost.localdomain.out
starting master, logging to /usr/local/hbase/logs/hbase-acadgild-master-localhost.localdomain.out
starting historyserver, logging to /usr/local/hadoop-2.6.0/logs/mapred-acadgild-historyserver-localhost.localdomain.out
Starting mysqld: [ OK ]
Starting Hive Metastore Server
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/hive-jdbc-0.14.0-standalone.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
```

- With the help of **jps** command we can list out the services that are currently in active state.

# Music Data Analysis

```
File Edit View Search Terminal Help
[acadgild@localhost ~]$
[acadgild@localhost ~]$ jps
3441 ResourceManager
5234 JobHistoryServer
5298 RunJar
2995 NameNode
3124 DataNode
3877 HMaster
3545 NodeManager
3292 SecondaryNameNode
5453 Jps
[acadgild@localhost ~]$ sudo service mysqld status
[sudo] password for acadgild:
mysqld (pid 5186) is running...
```

The start-daemon.sh script will check whether the current-batch.txt file is available in the logs folder or not. If not it will create the file and dump value '1' in that file and create LOGFILE with the current batchid.



## Lookup Tables creation in HBASE:

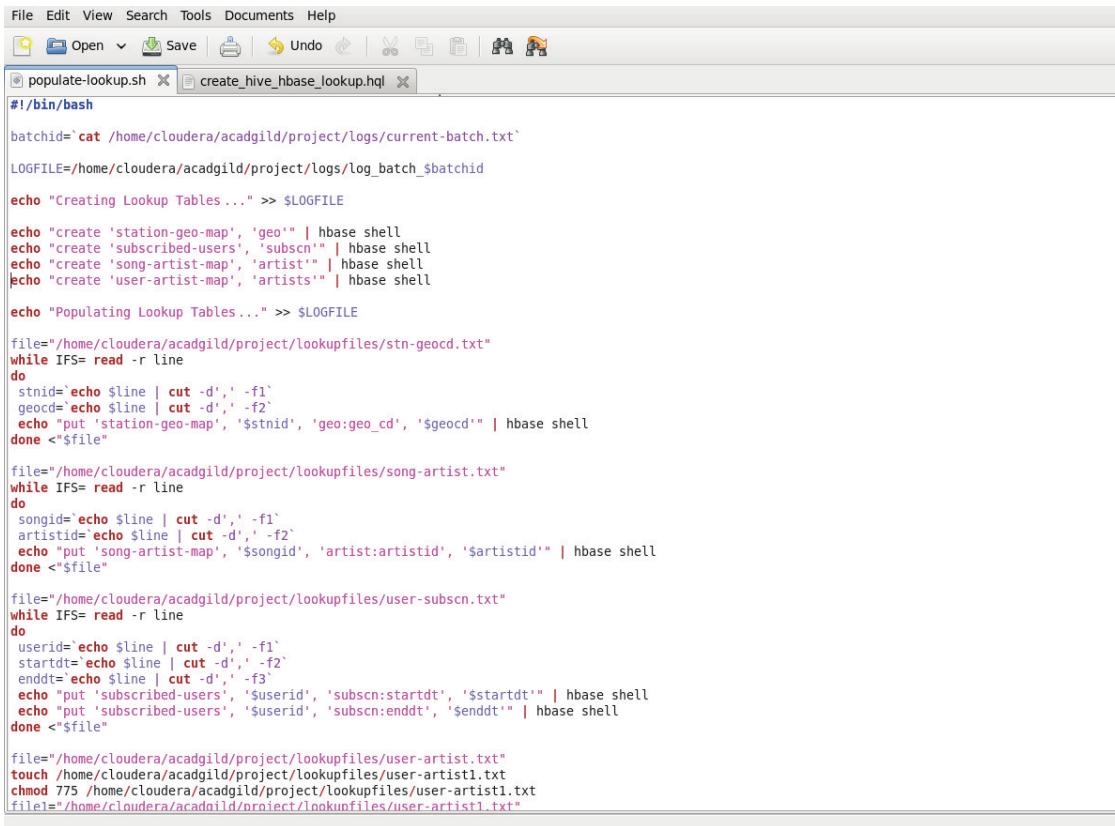
By using the populate-lookup.sh script we will create lookup tables in Hbase. These tables have to be used in data formatting, data enrichment and analysis stage.

Table Name	Description
Station_Geo_Map	Contains mapping of a geo_cd with station_id
Subscribed_Users	Contains user_id, subscription_start_date and subscription_end_date. Contains details only for subscribed users
Song_Artist_Map	Contains mapping of song_id with artist_id alongwith royalty associated with each play of the song
User_Artist_Map	Contains an array of artist_id(s) followed by a user_id

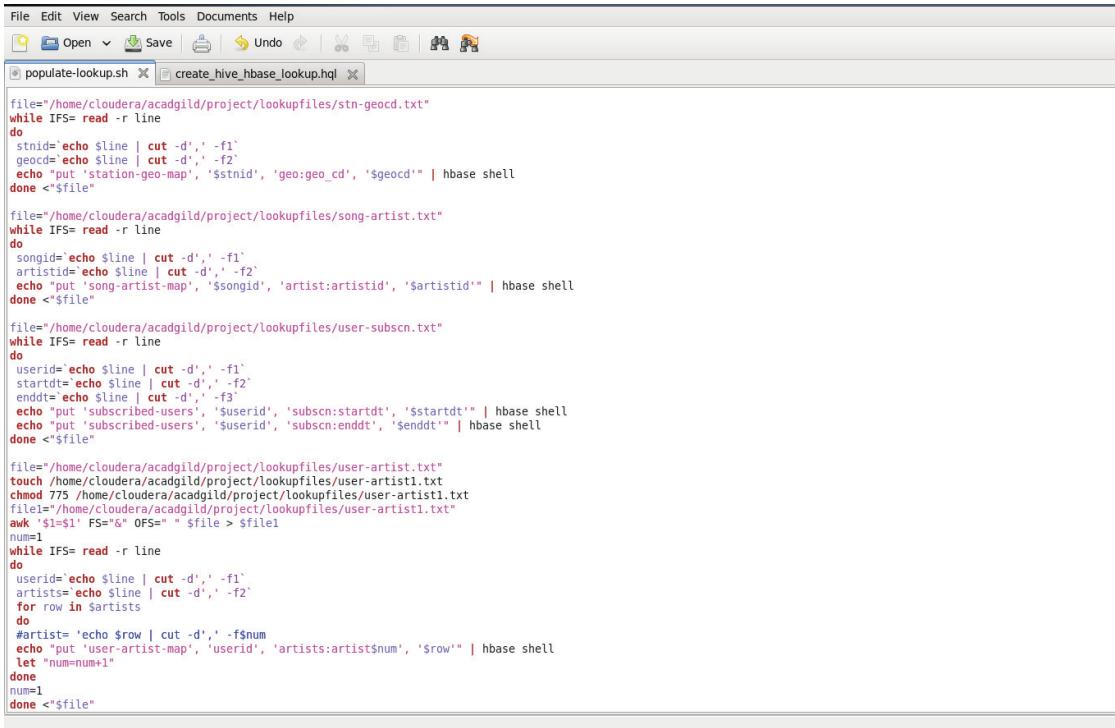
These 4 lookup files are used to create Hbase tables using populate-lookup.sh script and then with the help of data\_enrichment\_filtering\_schema.sh file we will create hive tables on the top of Hbase tables using create\_hive\_hbase\_lookup.hql Script.

# Music Data Analysis

## b) Creating Lookup Tables :



```
#!/bin/bash
batchid=`cat /home/cloudera/acadgild/project/logs/current-batch.txt`
LOGFILE=/home/cloudera/acadgild/project/logs/log_batch_$batchid
echo "Creating Lookup Tables ..." >> $LOGFILE
echo "create 'station-geo-map', 'geo'" | hbase shell
echo "create 'subscribed-users', 'subscn'" | hbase shell
echo "create 'song-artist-map', 'artist'" | hbase shell
echo "create 'user-artist-map', 'artists'" | hbase shell
echo "Populating Lookup Tables ..." >> $LOGFILE
file="/home/cloudera/acadgild/project/lookupfiles/stn-geocd.txt"
while IFS= read -r line
do
  snnid=`echo $line | cut -d',' -f1`
  geocd=`echo $line | cut -d',' -f2`
  echo "put 'station-geo-map', '$snnid', 'geo:geo_cd', '$geocd'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/song-artist.txt"
while IFS= read -r line
do
  songid=`echo $line | cut -d',' -f1`
  artistid=`echo $line | cut -d',' -f2`
  echo "put 'song-artist-map', '$songid', 'artist:artistid', '$artistid'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/user-subscn.txt"
while IFS= read -r line
do
  userid=`echo $line | cut -d',' -f1`
  startdt=`echo $line | cut -d',' -f2`
  enddt=`echo $line | cut -d',' -f3`
  echo "put 'subscribed-users', '$userid', 'subscn:startdt', '$startdt'" | hbase shell
  echo "put 'subscribed-users', '$userid', 'subscn:enddt', '$enddt'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/user-artist.txt"
touch /home/cloudera/acadgild/project/lookupfiles/user-artist1.txt
chmod 775 /home/cloudera/acadgild/project/lookupfiles/user-artist1.txt
file1="/home/cloudera/acadgild/project/lookupfiles/user-artist1.txt"
```



```
file="/home/cloudera/acadgild/project/lookupfiles/stn-geocd.txt"
while IFS= read -r line
do
  snnid=`echo $line | cut -d',' -f1`
  geocd=`echo $line | cut -d',' -f2`
  echo "put 'station-geo-map', '$snnid', 'geo:geo_cd', '$geocd'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/song-artist.txt"
while IFS= read -r line
do
  songid=`echo $line | cut -d',' -f1`
  artistid=`echo $line | cut -d',' -f2`
  echo "put 'song-artist-map', '$songid', 'artist:artistid', '$artistid'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/user-subscn.txt"
while IFS= read -r line
do
  userid=`echo $line | cut -d',' -f1`
  startdt=`echo $line | cut -d',' -f2`
  enddt=`echo $line | cut -d',' -f3`
  echo "put 'subscribed-users', '$userid', 'subscn:startdt', '$startdt'" | hbase shell
  echo "put 'subscribed-users', '$userid', 'subscn:enddt', '$enddt'" | hbase shell
done <"$file"
file="/home/cloudera/acadgild/project/lookupfiles/user-artist.txt"
touch /home/cloudera/acadgild/project/lookupfiles/user-artist1.txt
chmod 775 /home/cloudera/acadgild/project/lookupfiles/user-artist1.txt
file1="/home/cloudera/acadgild/project/lookupfiles/user-artist1.txt"
awk '$1=$1' FS="," OFS=" " $file > $file1
num=1
while IFS= read -r line
do
  userid=`echo $line | cut -d',' -f1`
  artists=`echo $line | cut -d',' -f2`
  for row in $artists
  do
    #Artist= `echo $row | cut -d',' -f$num
    echo "put 'user-artist-map', 'userid', 'artists:artist$num', '$row'" | hbase shell
    let "num=num+1"
  done
  num=1
done <"$file"
```

By using the shell scripting I created 4 lookup tables in Hbase NoSQL Database.

# Music Data Analysis

```
File Edit View Search Terminal Help
2017-08-31 20:28:01,616 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.98.14-hadoop2, r4e4aabb93b52f1b0fef6b66edd06ec8923014dec, Tue Aug 25 22:35:44 PDT 2015
'ut 'song-artist-map', 'S209', 'artist:artistid', 'A305'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-08-31 20:28:03,442 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
0 row(s) in 0.2850 seconds

2017-08-31 20:28:07,451 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.98.14-hadoop2, r4e4aabb93b52f1b0fef6b66edd06ec8923014dec, Tue Aug 25 22:35:44 PDT 2015
put 'subscribed-users', 'U100', 'subscn:startdt', '1465230523'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-08-31 20:28:09,229 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
0 row(s) in 0.2700 seconds

2017-08-31 20:28:13,179 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.98.14-hadoop2, r4e4aabb93b52f1b0fef6b66edd06ec8923014dec, Tue Aug 25 22:35:44 PDT 2015
'ut 'subscribed-users' 'U100' 'subscn:enddt' '1465130523'
```

```
File Edit View Search Terminal Help
2017-09-06 04:21:59,787 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
'ut 'user-artist-map','U112', 'artists:artist2','A301
0 row(s) in 0.2180 seconds

2017-09-06 04:22:06,838 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
put 'user-artist-map','U113', 'artists:artist1','A305'
0 row(s) in 0.2250 seconds

2017-09-06 04:22:14,019 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
'ut 'user-artist-map','U113', 'artists:artist2','A302
0 row(s) in 0.2180 seconds

2017-09-06 04:22:20,818 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
put 'user-artist-map','U114', 'artists:artist1','A300'
0 row(s) in 0.2260 seconds

2017-09-06 04:22:27,997 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
put 'user-artist-map','U114', 'artists:artist2','A301
0 row(s) in 0.2220 seconds

2017-09-06 04:22:35,159 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.10.0, rUnknown, Fri Jan 20 12:13:18 PST 2017
put 'user-artist-map','U114', 'artists:artist3','A302'
0 row(s) in 0.2270 seconds
```

Now we can see the lookup tables in Hbase shell terminal as shown below:

```
hbase(main):027:0> list
TABLE
song-artist-map
station-geo-map
subscribed-users
user-artist-map
4 row(s) in 0.0220 seconds

=> ["song-artist-map", "station-geo-map", "subscribed-users", "user-artist-map"]
hbase(main):028:0> █
```

# Music Data Analysis

```

File Edit View Search Terminal Help
hbase(main):023:0> scan 'user-artist-map'
ROW
U100
U100
U100
U101
U101
U102
U103
U103
U103
U104
U104
U105
U105
U105
U106
U106
U107
U108
U108
U108
U109
U109
U109
U110
U110
U111
U111
U112
U112
U113
U113
U114
U114
U114
15 row(s) in 0.1300 seconds

hbase(main):024:0> scan 'song-artist-map'
ROW
S200
S201
S202
S203
S204
S205
S206
S207
S208
S209
10 row(s) in 0.0530 seconds

```

```

File Edit View Search Terminal Help
hbase(main):025:0> scan 'station-geo-map'
ROW
ST400
ST401
ST402
ST403
ST404
ST405
ST406
ST407
ST408
ST409
ST410
ST411
ST412
ST413
ST414
15 row(s) in 0.0480 seconds

```

```

File Edit View Search Terminal Help
ST411
ST412
ST413
ST414
15 row(s) in 0.0480 seconds

```

```

hbase(main):026:0> scan 'subscribed-users'
ROW
U100
U100
U101
U101
U102
U102
U103
U103
U103
U104
U104
U105
U105
U105
U106
U106
U106
U107
U107
U107
U108
U108
U108
U109
U109
U109
U110
U110
U111
U111
U112
U112
U113
U113
U114
U114
U114
15 row(s) in 0.0590 seconds

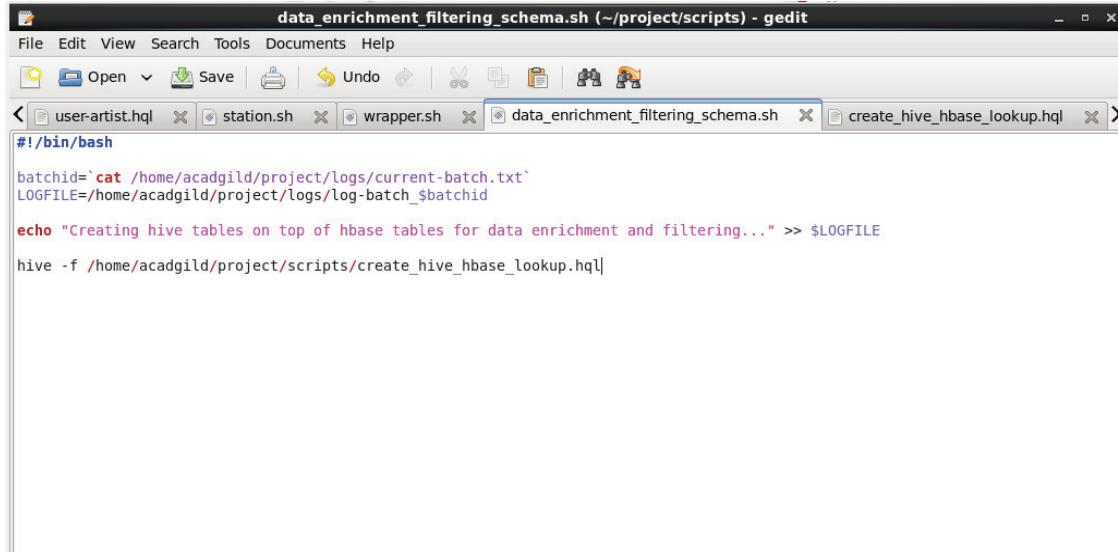
```

# Music Data Analysis

Now Lookup table creation is completed. So now we need to link these lookup tables in hive using the Hbase Storage Handler.

## c) Creating Hive Tables on the top of Hbase:

In this stage with the help of Hbase storage handler & SerDe properties we are creating the hive external tables by matching the columns of Hbase tables to hive tables.

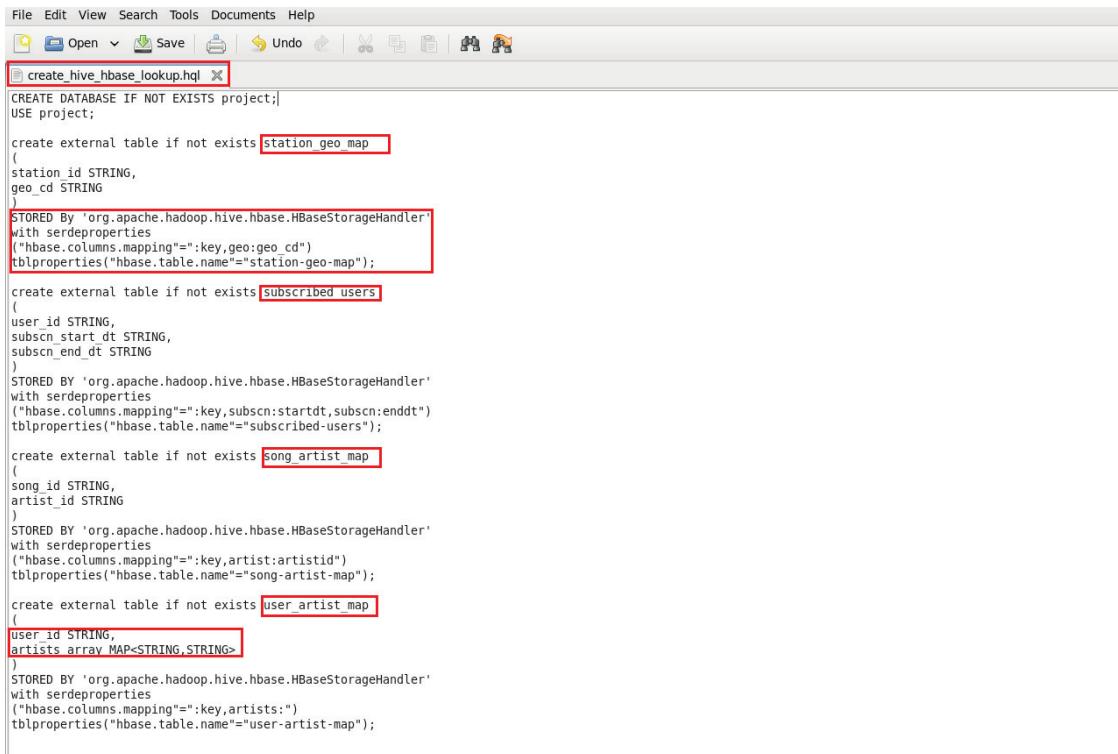


```
data_enrichment_filtering_schema.sh (~/project/scripts) - gedit
File Edit View Search Tools Documents Help
Open Save Undo
user-artist.hql station.sh wrapper.sh data_enrichment_filtering_schema.sh create_hive_hbase_lookup.hql
#!/bin/bash

batchid=`cat /home/acadgild/project/logs/current-batch.txt`
LOGFILE=/home/acadgild/project/logs/log-batch_$batchid

echo "Creating hive tables on top of hbase tables for data enrichment and filtering..." >> $LOGFILE

hive -f /home/acadgild/project/scripts/create_hive_hbase_lookup.hql
```



```
File Edit View Search Tools Documents Help
Open Save Undo
create_hive_hbase_lookup.hql

CREATE DATABASE IF NOT EXISTS project;
USE project;

create external table if not exists station_geo_map
(
station_id STRING,
geo_cd STRING
)
STORED By 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
with serdeproperties
("hbase.columns.mapping":":key,geo:geo cd")
tblproperties("hbase.table.name"="station-geo-map");

create external table if not exists subscribed_users
(
user_id STRING,
subscn_start_dt STRING,
subscn_end_dt STRING
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
with serdeproperties
("hbase.columns.mapping":":key,subscn:startdt,subscn:enddt")
tblproperties("hbase.table.name"="subscribed-users");

create external table if not exists song_artist_map
(
song_id STRING,
artist_id STRING
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
with serdeproperties
("hbase.columns.mapping":":key,artist:artistid")
tblproperties("hbase.table.name"="song-artist-map");

create external table if not exists user_artist_map
(
user_id STRING,
artists array MAP<STRING,STRING>
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
with serdeproperties
("hbase.columns.mapping":":key,artists:")
tblproperties("hbase.table.name"="user-artist-map");
```

- In the below screenshot we can see tables getting created in hive by running the data\_enrichement\_filtering\_schema.sh file.

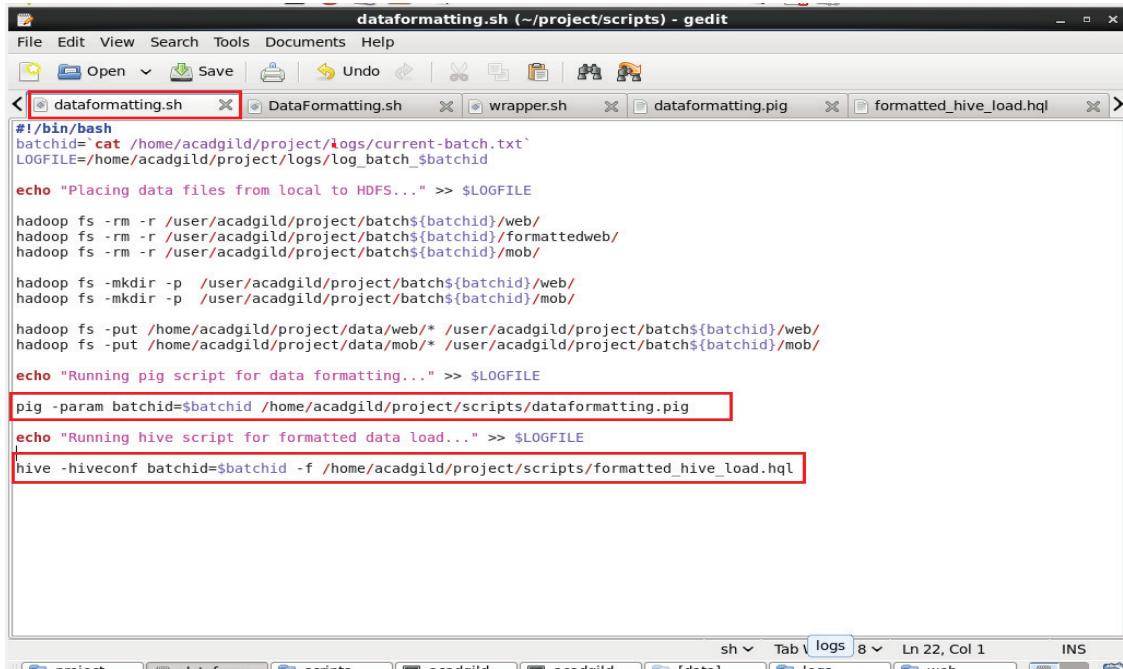
## Music Data Analysis

```
File Edit View Search Terminal Help
[cloudera@quickstart scripts]$ ./data_enrichment_filtering_schema.sh
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
OK
Time taken: 0.395 seconds
OK
Time taken: 0.102 seconds
OK
Time taken: 0.021 seconds
OK
Time taken: 0.018 seconds
OK
Time taken: 1.085 seconds
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[cloudera@quickstart scripts]$
```

```
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> show databases;
OK
default
project
Time taken: 0.499 seconds, Fetched: 2 row(s)
hive> use project;
OK
Time taken: 0.047 seconds
hive> show tables;
OK
song_artist_map
station_geo_map
subscribed_users
user_artist_map
Time taken: 0.049 seconds, Fetched: 4 row(s)
hive> ■
```

### d) Data Formatting:

In this stage we are merging the data coming from both web applications and mobile applications and create a common table for analyzing purpose and create partitioned data based on batchid, since we are running this scripts for every 3 hours.



```
dataformatting.sh (~/project/scripts) - gedit
File Edit View Search Tools Documents Help
File Open Save Undo
dataformatting.sh wrapper.sh dataformatting.pig formatted_hive_load.hql
#!/bin/bash
batchid= cat /home/acadgild/project/logs/current-batch.txt
LOGFILE=/home/acadgild/project/logs/log_batch_$batchid
echo "placing data files from local to HDFS..." >> $LOGFILE
hadoop fs -rm -r /user/acadgild/project/batch${batchid}/web/
hadoop fs -rm -r /user/acadgild/project/batch${batchid}/formattedweb/
hadoop fs -rm -r /user/acadgild/project/batch${batchid}/mob/
hadoop fs -mkdir -p /user/acadgild/project/batch${batchid}/web/
hadoop fs -mkdir -p /user/acadgild/project/batch${batchid}/mob/
hadoop fs -put /home/acadgild/project/data/web/* /user/acadgild/project/batch${batchid}/web/
hadoop fs -put /home/acadgild/project/data/mob/* /user/acadgild/project/batch${batchid}/mob/
echo "Running pig script for data formatting..." >> $LOGFILE
pig -param batchid=$batchid /home/acadgild/project/scripts/dataformatting.pig
echo "Running hive script for formatted data load..." >> $LOGFILE
hive -hiveconf batchid=$batchid -f /home/acadgild/project/scripts/formatted_hive_load.hql
```

I am running two scripts to format the data. They are: 1) Dataformatting.pig

2) Formatted\_hive\_load.hql

Pig script to parse the data from coming from web\_data.xml to csv format and partition both web and mob data based on based on batch ID's



# Music Data Analysis

```
File Edit View Search Terminal Help
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
[cloudera@quickstart scripts]$ ./dataformatting.sh
Deleted /user/acadgild/project/batch1/web
`Lrm: '/user/acadgild/project/batch1/formattedweb/: No such file or directory
Deleted /user/acadgild/project/batch1/mob
log4j:WARN appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
2017-09-06 04:41:09:677 [main] INFO org.apache.pig.Main - Apache Pig version 0.12.0-cdh5.10.0 (reported) compiled Jan 29 2017, 12:05:43
2017-09-06 04:41:09:678 [main] INFO org.apache.pig.Main - Logging error messages to: /home/cloudera/acadgild/project/scripts/pig_1504698867668.log
2017-09-06 04:41:09:838 [main] INFO org.apache.pig.impl.util.Utils - Default bootstrap file /home/cloudera/.pigbootup not found
2017-09-06 04:41:09:849 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:950 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:950 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: hdfs://quickstart.cloudera:8020
2017-09-06 04:41:09:951 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to map-reduce job tracker at: localhost:8021
2017-09-06 04:41:09:954 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:955 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:956 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:957 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:958 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:959 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:960 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:961 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:962 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:963 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:968 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:969 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-06 04:41:09:970 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:973 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:978 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:979 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:980 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:981 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:982 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:983 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:984 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:985 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:09:986 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-06 04:41:10:498 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupE
timerizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushdownForEachFlatten, PushupFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[Filter
]}
2017-09-06 04:41:10:587 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.s
2017-09-06 04:41:10:694 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2017-09-06 04:41:10:725 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - MR plan size before optimization: 1
2017-09-06 04:41:10:725 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - MR plan size after optimization: 1
2017-09-06 04:41:10:725 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - MR plan size after optimization: 1
2017-09-06 04:41:11:032 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - mapred.job.reduce.sets is deprecated. Instead, use mapreduce.job.reduces
2017-09-06 04:41:11:074 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2017-09-06 04:41:11:176 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.reduce.markreset.buffer.percent is deprecated. Instead, use mapreduce.reduce.markreset
2017-09-06 04:41:11:176 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to defa
2017-09-06 04:41:11:176 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.output.compress is deprecated. Instead, use mapreduce.output.fileoutputformat.compress
2017-09-06 04:41:11:838 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job387487764963206215.jar
2017-09-06 04:41:15:855 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job387487764963206215.jar created
2017-09-06 04:41:15:855 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.deprecated. Instead, use mapreduce.job.jar
2017-09-06 04:41:15:876 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up cloudera@quickstart:-
```

```
File Edit View Search Terminal Help
2017-09-05 23:29:02,615 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2017-09-05 23:29:16,031 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2017-09-05 23:29:17,880 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2017-09-05 23:29:17,850 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2017-09-05 23:29:17,851 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:
Success!
Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReducetime Alias Feature C
job_1504678625866_0001 1 0 4 4 4 n/a n/a n/a A,B MAP_ONLY /user/acadgild/project/batch1/formattedweb,
Input(s):
Successfully read 20 records (7353 bytes) from: "/user/acadgild/project/batch1/web"
Output(s):
Successfully stored 20 records (1236 bytes) in: "/user/acadgild/project/batch1/formattedweb"
Counters:
Total records written : 20
Total bytes written : 1236
Spillable Memory Manager spill count : 0
Total bags proactively spilled : 0
Total records proactively spilled : 0
Job DAG:
job_1504678625866_0001

2017-09-05 23:29:17,930 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
DK
Time taken: 0.393 seconds
DK
Time taken: 0.446 seconds
Loading data to table project.formatted input partition (batchid=1)
Partition project.formatted_input{batchid=1} stats: [numFiles=1, numRows=0, totalSize=1236, rawDataSize=0]
DK
Time taken: 0.908 seconds
Loading data to table project.formatted_input partition (batchid=1)
Partition project.formatted_input{batchid=1} stats: [numFiles=2, numRows=0, totalSize=2493, rawDataSize=0]
DK
Time taken: 0.742 seconds
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
cloudera@quickstart:~
```

In the above screenshot we can see the dataformatting.pig along with the formatted\_hive\_load.hql executed successfully.



# Music Data Analysis

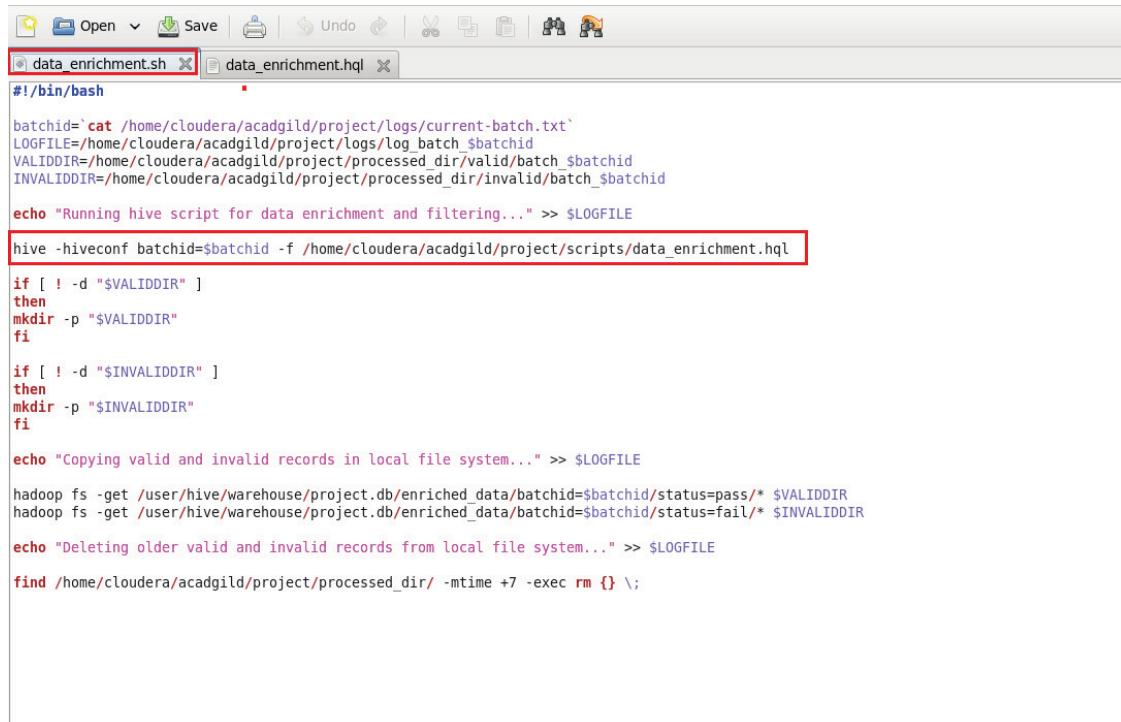
## e) Data Enrichment:

In this phase we will enrich the data coming from web and mobile applications using the lookup table stored in Hbase and divide the records based on the enrichment rules into 'pass' and 'fail' records.

Rules for data enrichment

1. If any of like or dislike is NULL or absent, consider it as 0.
2. If fields like Geo\_cd and Artist\_id are NULL or absent, consult the lookup tables for fields Station\_id and Song\_id respectively to get the values of Geo\_cd and Artist\_id.
3. If corresponding lookup entry is not found, consider that record to be invalid

So based on the enrichment rules we will fill the null geo\_cd and artist\_id values with the help of corresponding lookup values in song-artist-map and station-geo-map tables in Hive-Hbase tables.



```
#!/bin/bash
batchid=`cat /home/cloudera/acadgild/project/logs/current-batch.txt`
LOGFILE=/home/cloudera/acadgild/project/logs/log_batch_$batchid
VALIDDIR=/home/cloudera/acadgild/project/processed_dir/valid/batch_$batchid
INVALIDDIR=/home/cloudera/acadgild/project/processed_dir/invalid/batch_$batchid

echo "Running hive script for data enrichment and filtering..." >> $LOGFILE
hive -hiveconf batchid=$batchid -f /home/cloudera/acadgild/project/scripts/data_enrichment.hql

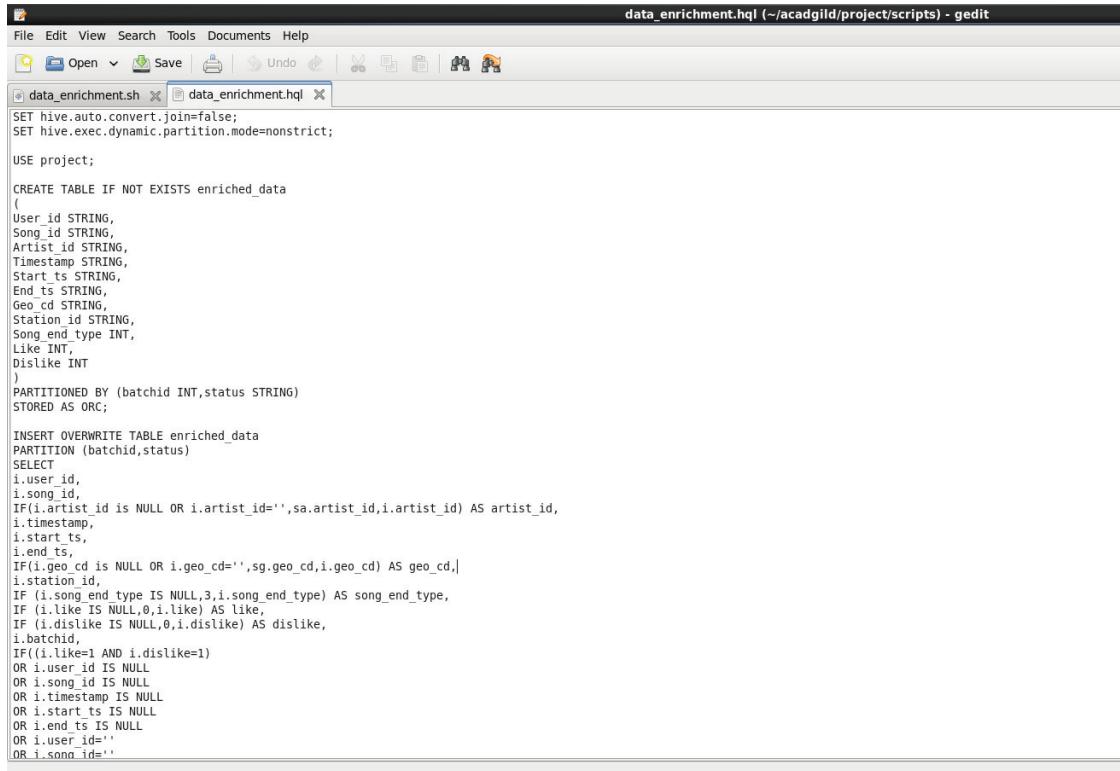
if [ ! -d "$VALIDDIR" ]
then
mkdir -p "$VALIDDIR"
fi

if [ ! -d "$INVALIDDIR" ]
then
mkdir -p "$INVALIDDIR"
fi

echo "Copying valid and invalid records in local file system..." >> $LOGFILE
hadoop fs -get /user/hive/warehouse/project.db/enriched_data/batchid=$batchid/status=pass/* $VALIDDIR
hadoop fs -get /user/hive/warehouse/project.db/enriched_data/batchid=$batchid/status=fail/* $INVALIDDIR

echo "Deleting older valid and invalid records from local file system..." >> $LOGFILE
find /home/cloudera/acadgild/project/processed_dir/ -mtime +7 -exec rm {} \;
```

# Music Data Analysis

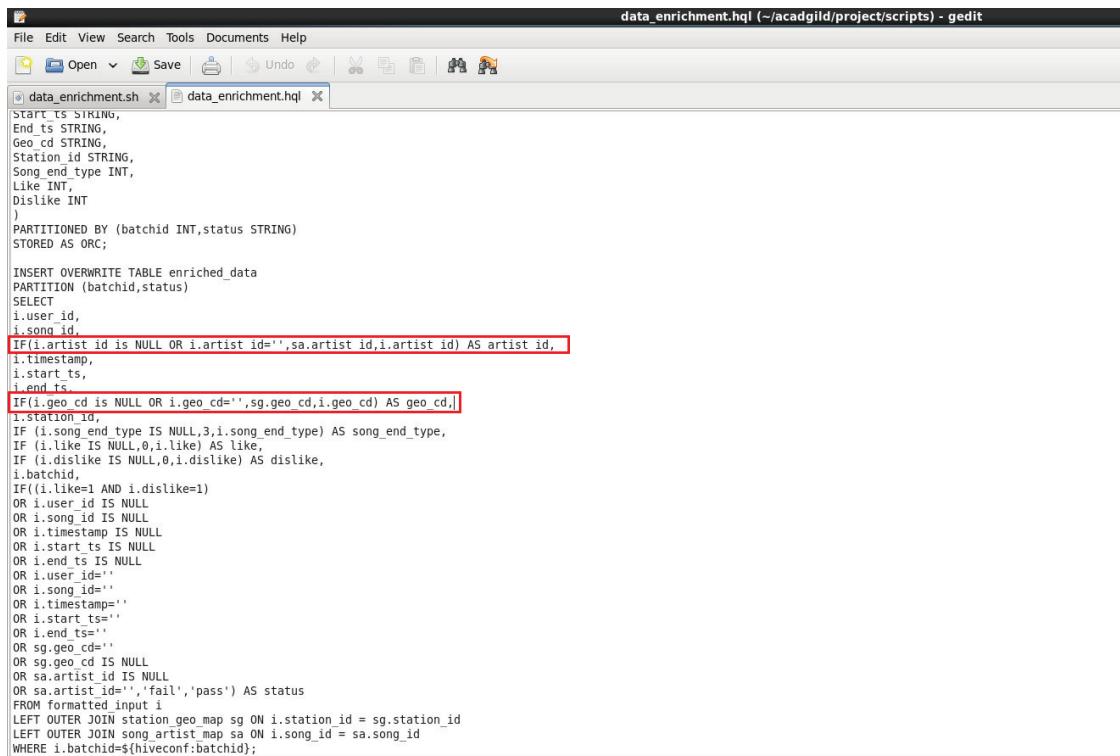


```
data_enrichment.hql (~/acadgild/project/scripts) - gedit
File Edit View Search Tools Documents Help
data_enrichment.sh data_enrichment.hql
SET hive.auto.convert.join=false;
SET hive.exec.dynamic.partition.mode=nonstrict;

USE project;

CREATE TABLE IF NOT EXISTS enriched_data
(
User_id STRING,
Song_id STRING,
Artist_id STRING,
Timestamp STRING,
Start_ts STRING,
End_ts STRING,
Geo_cd STRING,
Station_id STRING,
Song_end_type INT,
Like INT,
Dislike INT
)
PARTITIONED BY (batchid INT,status STRING)
STORED AS ORC;

INSERT OVERWRITE TABLE enriched_data
PARTITION (batchid,status)
SELECT
i.user_id,
i.song_id,
IF(i.artist_id is NULL OR i.artist_id='',sa.artist_id,i.artist_id) AS artist_id,
i.timestamp,
i.start_ts,
i.end_ts,
IF(i.geo_cd is NULL OR i.geo_cd='',sg.geo_cd,i.geo_cd) AS geo_cd,
i.station_id,
IF (i.song_end_type IS NULL,3,i.song_end_type) AS song_end_type,
IF (i.like IS NULL,0,i.like) AS like,
IF (i.dislike IS NULL,0,i.dislike) AS dislike,
i.batchid,
IF(i.like=1 AND i.dislike=1)
OR i.user_id IS NULL
OR i.song_id IS NULL
OR i.timestamp IS NULL
OR i.start_ts IS NULL
OR i.end_ts IS NULL
OR i.user_id=''
OR i.song_id=''
IF(i.artist_id is NULL OR i.artist_id='',sa.artist_id,i.artist_id) AS artist_id,
i.timestamp,
i.start_ts,
i.end_ts,
IF(i.geo_cd is NULL OR i.geo_cd='',sg.geo_cd,i.geo_cd) AS geo_cd,
i.station_id,
IF (i.song_end_type IS NULL,3,i.song_end_type) AS song_end_type,
IF (i.like IS NULL,0,i.like) AS like,
IF (i.dislike IS NULL,0,i.dislike) AS dislike,
i.batchid,
IF(i.like=1 AND i.dislike=1)
OR i.user_id IS NULL
OR i.song_id IS NULL
OR i.timestamp IS NULL
OR i.start_ts IS NULL
OR i.end_ts IS NULL
OR i.user_id=''
OR i.song_id=''
OR i.timestamp=''
OR i.start_ts=''
OR i.end_ts=''
OR sg.geo_cd=''
OR sg.geo_cd IS NULL
OR sa.artist_id IS NULL
OR sa.artist_id='', 'fail', 'pass') AS status
FROM formatted_input i
LEFT OUTER JOIN station geo map sg ON i.station_id = sg.station_id
LEFT OUTER JOIN song_artist map sa ON i.song_id = sa.song_id
WHERE i.batchid=${hiveconf:batchid};
```



```
data_enrichment.hql (~/acadgild/project/scripts) - gedit
File Edit View Search Tools Documents Help
data_enrichment.sh data_enrichment.hql
Start_ts STRING,
End_ts STRING,
Geo_cd STRING,
Station_id STRING,
Song_end_type INT,
Like INT,
Dislike INT
)
PARTITIONED BY (batchid INT,status STRING)
STORED AS ORC;

INSERT OVERWRITE TABLE enriched_data
PARTITION (batchid,status)
SELECT
i.user_id,
i.song_id,
IF(i.artist_id is NULL OR i.artist_id='',sa.artist_id,i.artist_id) AS artist_id,
i.timestamp,
i.start_ts,
i.end_ts,
IF(i.geo_cd is NULL OR i.geo_cd='',sg.geo_cd,i.geo_cd) AS geo_cd,
i.station_id,
IF (i.song_end_type IS NULL,3,i.song_end_type) AS song_end_type,
IF (i.like IS NULL,0,i.like) AS like,
IF (i.dislike IS NULL,0,i.dislike) AS dislike,
i.batchid,
IF(i.like=1 AND i.dislike=1)
OR i.user_id IS NULL
OR i.song_id IS NULL
OR i.timestamp IS NULL
OR i.start_ts IS NULL
OR i.end_ts IS NULL
OR i.user_id=''
OR i.song_id=''
OR i.timestamp=''
OR i.start_ts=''
OR i.end_ts=''
OR sg.geo_cd=''
OR sg.geo_cd IS NULL
OR sa.artist_id IS NULL
OR sa.artist_id='', 'fail', 'pass') AS status
FROM formatted_input i
LEFT OUTER JOIN station geo map sg ON i.station_id = sg.station_id
LEFT OUTER JOIN song_artist map sa ON i.song_id = sa.song_id
WHERE i.batchid=${hiveconf:batchid};
```

# Music Data Analysis

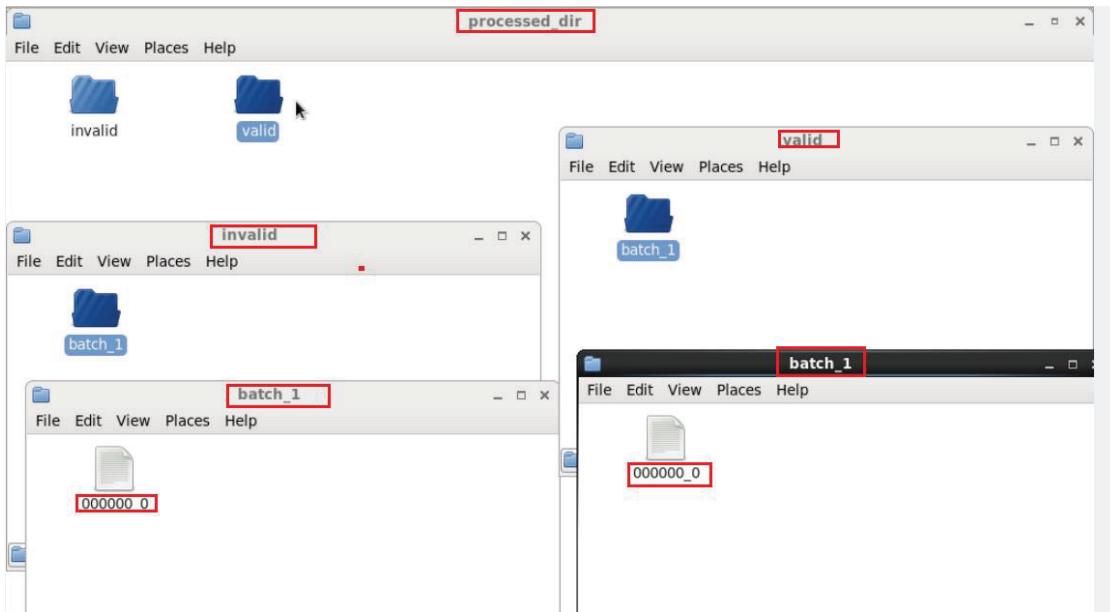
```
File Edit View Search Terminal Help
[cloudera@quickstart project]$ cd scripts/
[cloudera@quickstart scripts]$ ./data_enrichment.sh

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
OK
Time taken: 0.391 seconds
OK
Time taken: 0.178 seconds
Query ID = cloudera_20170905233838_f151d27a-601b-405e-9313-6d9df9bd18f2
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1504678625866_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1504678625866_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1504678625866_0002
Hadoop job information for Stage-1: number of mappers: 3; number of reducers: 1
2017-09-05 23:38:53,407 Stage-1 map = 0%, reduce = 0%
2017-09-05 23:39:03,399 Stage-1 map = 33%, reduce = 0%, Cumulative CPU 1.06 sec
2017-09-05 23:39:04,436 Stage-1 map = 67%, reduce = 0%, Cumulative CPU 2.12 sec
2017-09-05 23:39:05,492 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.05 sec
2017-09-05 23:39:10,763 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.5 sec
MapReduce Total cumulative CPU time: 5 seconds 500 msec
Ended Job = job_1504678625866_0002
Launching Job 2 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1504678625866_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1504678625866_0003/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1504678625866_0003
Hadoop job information for Stage-2: number of mappers: 2; number of reducers: 1
2017-09-05 23:39:25,758 Stage-2 map = 0%, reduce = 0%
2017-09-05 23:39:33,225 Stage-2 map = 50%, reduce = 0%, Cumulative CPU 0.95 sec
2017-09-05 23:39:34,284 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.85 sec
2017-09-05 23:39:40,558 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.58 sec
MapReduce Total cumulative CPU time: 4 seconds 580 msec
Ended Job = job_1504678625866_0003
Loading data to table project.enriched_data partition (batchid=null, status=null)
  Time taken for load dynamic partitions : 222
    Loading partition {batchid=1, status=fail}
    Loading partition {batchid=1, status=pass}
  Time taken for adding to write entity : 2
```

```
File Edit View Search Terminal Help
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1504678625866_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1504678625866_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1504678625866_0002
Hadoop job information for Stage-1: number of mappers: 3; number of reducers: 1
2017-09-05 23:38:53,407 Stage-1 map = 0%, reduce = 0%
2017-09-05 23:39:03,399 Stage-1 map = 33%, reduce = 0%, Cumulative CPU 1.06 sec
2017-09-05 23:39:04,436 Stage-1 map = 67%, reduce = 0%, Cumulative CPU 2.12 sec
2017-09-05 23:39:05,492 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.05 sec
2017-09-05 23:39:10,763 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.5 sec
MapReduce Total cumulative CPU time: 5 seconds 500 msec
Ended Job = job_1504678625866_0002
Launching Job 2 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1504678625866_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1504678625866_0003/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1504678625866_0003
Hadoop job information for Stage-2: number of mappers: 2; number of reducers: 1
2017-09-05 23:39:25,758 Stage-2 map = 0%, reduce = 0%
2017-09-05 23:39:33,225 Stage-2 map = 50%, reduce = 0%, Cumulative CPU 0.95 sec
2017-09-05 23:39:34,284 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.85 sec
2017-09-05 23:39:40,558 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.58 sec
MapReduce Total cumulative CPU time: 4 seconds 580 msec
Ended Job = job_1504678625866_0003
Loading data to table project.enriched_data partition (batchid=null, status=null)
  Time taken for load dynamic partitions : 222
    Loading partition {batchid=1, status=fail}
    Loading partition {batchid=1, status=pass}
  Time taken for adding to write entity : 2
Partition project.enriched_data{batchid=1, status=fail} stats: [numFiles=1, numRows=18, totalSize=1481, rawDataSize=12435]
Partition project.enriched_data{batchid=1, status=pass} stats: [numFiles=1, numRows=22, totalSize=1487, rawDataSize=16126]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 3 Reduce: 1 Cumulative CPU: 5.5 sec HDFS Read: 45891 HDFS Write: 3097 SUCCESS
Stage-Stage-2: Map: 2 Reduce: 1 Cumulative CPU: 4.58 sec HDFS Read: 21135 HDFS Write: 3135 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 80 msec
OK
Time taken: 63.499 seconds
WARN: The method class org.apache.commons.logging.impl.SLF4JLogFactory#release() was invoked.
WARN: Please see http://www.slf4j.org/codes.html#release for an explanation.
```

At the end script will automatically divide the records based on status pass & fail and dump the result into processed\_dir folder with valid and invalid folders.

## Music Data Analysis



- Now we can check whether the data properly loaded in the hive terminal or not .

```
hive> show tables;
OK
enriched_data
formatted_input
song_artist_map
station_geo_map
subscribed_users
users_artists
Time taken: 0.034 seconds, Fetched: 6 row(s)
```

- In the below screenshot we have data for data enrichment table where we filled the null values of artist\_id and geo\_cd of formatted input with the help of lookup tables

# Music Data Analysis

```

File Edit View Search Terminal Help
formatted input
song_artist_map
station_geo_map
subscribed_users
unsubscribed_users
user_artist_map
Time taken: 0.214 seconds, Fetched: 7 row(s)
hive> select * from enriched data;
OK
U113 S200 A303 1465230523 1475130523 1465130523 E ST413 3 1 1 1 1 fail
U100 S200 A301 1494342562 1494342562 AP ST410 3 1 1 1 1 fail
U107 S202 A303 1495130523 1465230523 U ST415 0 1 1 1 1 fail
U103 S202 A300 1465535556 1465535556 A ST415 2 1 1 1 1 fail
U106 S202 A302 1465230523 1465130523 1465130523 AU ST408 0 1 1 1 1 fail
U109 S203 A304 1462908262 1494342562 1468139889 E ST405 1 1 1 1 1 fail
S203 A302 1495130523 1475130523 1465230523 E ST400 0 0 1 1 1 fail
U110 S203 A300 1465230523 1485130523 A ST415 0 1 1 1 1 fail
U111 S204 A300 1465535556 1468139889 U ST410 3 1 1 1 1 fail
U113 S204 A301 1494342562 1494342562 1465535556 E ST415 3 0 1 1 1 fail
U100 S204 A302 1495130523 1475130523 1465130523 AU ST408 2 1 1 1 1 fail
U106 S205 A300 1462908262 1462908262 1494342562 AP ST407 2 1 1 1 1 fail
U111 S206 A305 1465130523 1485130523 AU ST415 0 1 1 1 1 fail
U114 S207 A303 1465130523 1465230523 1475130523 A ST415 3 1 0 1 1 fail
U102 S207 A301 1465230523 1485130523 1465230523 AU ST403 3 1 1 1 1 fail
S208 A300 1465535556 1494342562 1465535556 E ST411 1 0 1 1 1 fail
U118 S208 A303 1475130523 1465130523 1465230523 E ST415 3 0 0 1 1 fail
U119 S208 A302 1495130523 1465230523 1465230523 U ST415 3 0 0 0 1 fail
U101 S208 A300 1462908262 1468139889 1462908262 E ST408 0 1 1 1 1 fail
U107 S218 A302 1475130523 1485130523 1485130523 AP ST404 2 1 0 1 1 fail
U115 S206 A306 1465535556 1494342562 1465535556 AU ST404 3 0 0 1 1 pass
U108 S209 A302 1468139889 1462908262 1468139889 U ST414 0 0 1 1 1 pass
U120 S201 A300 1494342562 1465535556 1468139889 A ST410 3 0 0 1 1 pass
U107 S202 A304 1494342562 1468139889 1462908262 U ST409 0 0 0 1 1 pass
U101 S202 A300 1465230523 1465130523 1475130523 U ST401 0 0 1 1 1 pass
U110 S202 A303 1494342562 1494342562 1468139889 AU ST402 2 1 0 1 1 pass
U118 S202 A300 1495130523 1465230523 1465230523 AP ST410 1 0 0 0 1 pass
U104 S202 A303 1465230523 1475130523 1465130523 A ST409 2 0 0 1 1 pass
U102 S203 A305 1465535556 1465535556 1494342562 U ST404 2 0 0 0 1 pass
U113 S203 A304 1465535556 1462908262 1465230523 U ST405 0 0 1 1 1 pass
U113 S203 A303 1465535556 1468139889 1494342562 U ST408 2 0 0 0 1 pass
U105 S203 A303 1475130523 1465230523 1465130523 AU ST408 2 0 0 1 1 pass
U113 S203 A303 1465535556 1465535556 A ST402 1 1 0 0 1 pass
U103 S204 A300 1468139889 1494342562 1465535556 AU ST411 2 1 0 0 1 pass
U100 S205 A304 1465130523 1475130523 1465130523 A ST410 2 1 0 0 1 pass
U106 S206 A300 1494342562 1465535556 1462908262 A ST405 3 1 0 0 1 pass
U120 S206 A303 1495130523 1485130523 1465130523 AU ST414 0 0 0 1 1 pass
U105 S207 A300 1465230523 1485130523 1465130523 U ST400 2 0 0 1 1 pass
U116 S208 A303 1465230523 1485130523 1475130523 A ST413 1 0 0 1 1 pass
U114 S209 A303 1465535556 1462908262 1494342562 U ST411 2 1 0 0 1 pass
Time taken: 0.419 seconds, Fetched: 40 row(s)

```

- Enrichment phase is executed successfully by applying all the rules of enrichment.

## f) Data Analysis using Spark:

- In this stage we will do analysis on enriched data using Spark SQL and run the program using **Spark-Submit** command.
- Before running the spark-submit command we have to **zip -d** command to remove the bad manifests in created spark project jar file to avoid the invalid Signature exception.
- I used two spark-submits for analysis. 1) Spark\_analysis for creating tables for each query/problem statement.  
2) Spark\_analysis\_2 for displaying results for each query in terminal.

```

#!/bin/bash
batchid=`cat /home/acadgild/project/logs/current-batch.txt`

LOGFILE=/home/acadgild/project/logs/log_batch_$batchid

echo "Running script for data analysis using spark..." >> $LOGFILE
chmod 775 /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar

zip -d /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar META-INF/*.DSA META-INF/*.RSA META-INF/*.SF

spark-submit \
--class Spark_analysis \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid

spark-submit \
--class Spark_analysis_2 \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid

echo "Exporting data to MYSQL using sqoop export..." >> $LOGFILE
sh /home/acadgild/project/scripts/data_export.sh

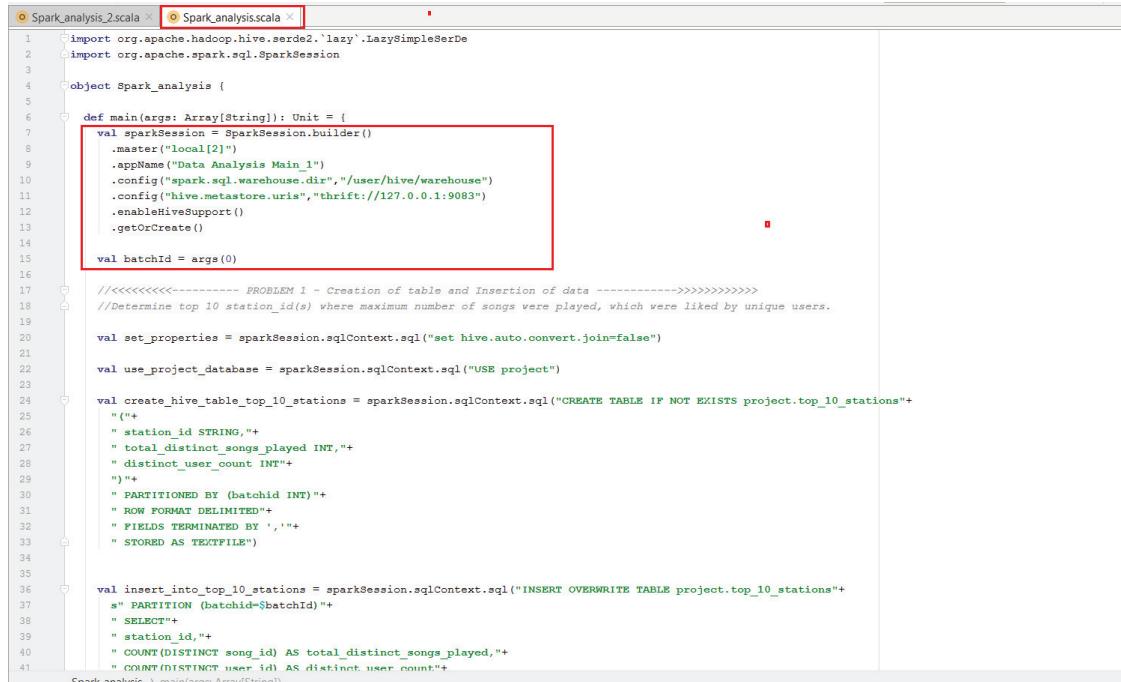
echo "Incrementing batchid..." >> $LOGFILE
batchid=`expr $batchid + 1`
echo -n $batchid > /home/acadgild/project/logs/current-batch.txt

```

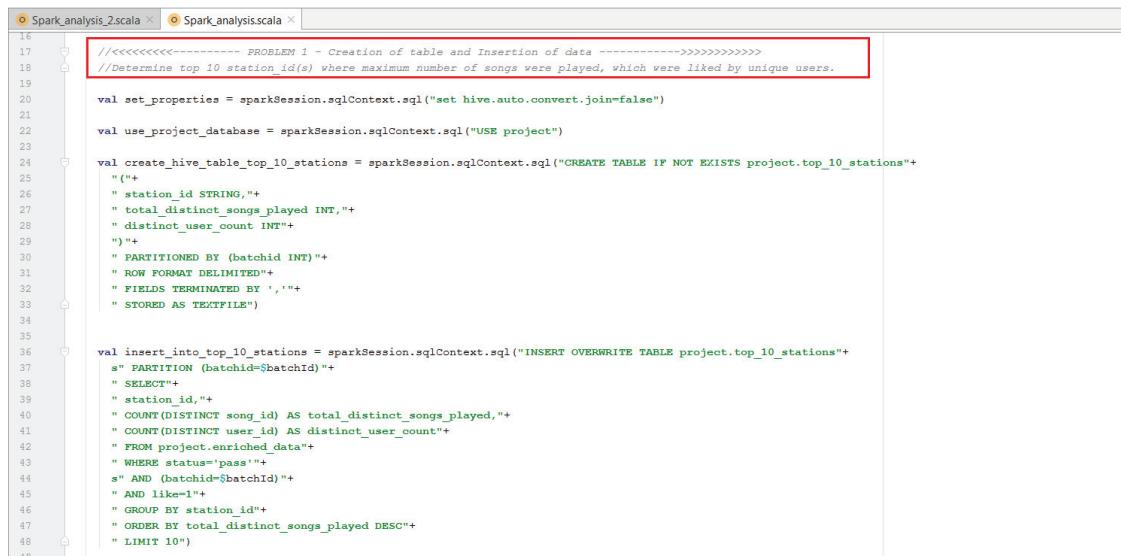
# Music Data Analysis

## Spark Source Code:

I have created one Scala file for creating tables as per query wise.



```
1 import org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
2 import org.apache.spark.sql.SparkSession
3
4 object Spark_analysis {
5
6   def main(args: Array[String]): Unit = {
7     val sparkSession = SparkSession.builder()
8       .master("local[2]")
9       .appName("Data Analysis Main_1")
10      .config("spark.sql.warehouse.dir","/user/hive/warehouse")
11      .config("hive.metastore.uris","thrift://127.0.0.1:9083")
12      .enableHiveSupport()
13      .getOrCreate()
14
15     val batchId = args(0)
16
17     //<<<<<<----- PROBLEM 1 - Creation of table and Insertion of data ----->>>>>>
18     //Determine top 10 station_id(s) where maximum number of songs were played, which were liked by unique users.
19
20     val set_properties = sparkSession.sqlContext.sql("set hive.auto.convert.join=false")
21
22     val use_project_database = sparkSession.sqlContext.sql("USE project")
23
24     val create_hive_table_top_10_stations = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.top_10_stations"+
25       "(*+
26       " station_id STRING,"+
27       " total_distinct_songs_played INT,"+
28       " distinct_user_count INT"+"
29     )"+"
30     " PARTITIONED BY (batchid INT)" +
31     " ROW FORMAT DELIMITED"+"
32     " FIELDS TERMINATED BY ','"+"
33     " STORED AS TEXTFILE")
34
35
36     val insert_into_top_10_stations = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.top_10_stations"+
37       "s" PARTITION (batchid=$batchId)" +
38       " SELECT"+"
39       " station_id,"+
40       " COUNT(DISTINCT song_id) AS total_distinct_songs_played,"+
41       " COUNT(DISTINCT user_id) AS distinct_user_count"+"
42     FROM project.enriched_data"+
43     " WHERE status='pass'"+
44     " AND (batchid=$batchId)" +
45     " AND like=1" +
46     " GROUP BY station_id" +
47     " ORDER BY total_distinct_songs_played DESC" +
48     " LIMIT 10")
49
50   }
```



```
16 //<<<<<<----- PROBLEM 1 - Creation of table and Insertion of data ----->>>>>>
17 //Determine top 10 station_id(s) where maximum number of songs were played, which were liked by unique users.
18
19 val set_properties = sparkSession.sqlContext.sql("set hive.auto.convert.join=false")
20
21 val use_project_database = sparkSession.sqlContext.sql("USE project")
22
23 val create_hive_table_top_10_stations = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.top_10_stations"+
24   "(*+
25   " station_id STRING,"+
26   " total_distinct_songs_played INT,"+
27   " distinct_user_count INT"+"
28 )"+"
29   " PARTITIONED BY (batchid INT)" +
30   " ROW FORMAT DELIMITED"+"
31   " FIELDS TERMINATED BY ','"+"
32   " STORED AS TEXTFILE")
33
34
35 val insert_into_top_10_stations = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.top_10_stations"+
36   "s" PARTITION (batchid=$batchId)" +
37   " SELECT"+"
38   " station_id,"+
39   " COUNT(DISTINCT song_id) AS total_distinct_songs_played,"+
40   " COUNT(DISTINCT user_id) AS distinct_user_count"+"
41   " FROM project.enriched_data"+
42   " WHERE status='pass'"+
43   " AND (batchid=$batchId)" +
44   " AND like=1" +
45   " GROUP BY station_id" +
46   " ORDER BY total_distinct_songs_played DESC" +
47   " LIMIT 10")
```

# Music Data Analysis

```

49
50 //<<<<<<----- PROBLEM 2 - Creation of table and Insertion of data ----->>>>>>>>
51 /*Determine total duration of songs played by each type of user, where type of user can be 'subscribed' or 'unsubscribed'.
52 An unsubscribed user is the one whose record is either not present in Subscribed_users lookup table or has subscription_end_date
53 earlier than the timestamp of the song played by him.*/
54
55 val create_hive_table_song_duration = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.song_duration"+
56 " ("+
57 " user_id STRING,"+
58 " user_type STRING,"+
59 " song_id STRING,"+
60 " artist_id STRING,"+
61 " total_duration_in_minutes DOUBLE"+
62 ")"+
63 " PARTITIONED BY (batchid INT)"+
64 " ROW FORMAT DELIMITED"+
65 " FIELDS TERMINATED BY ','"+
66 " STORED AS TEXTFILE")
67
68
69 val insert_into_song_duration = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.song_duration"+
70 " PARTITION (batchid=$batchId)"+
71 " SELECT"+
72 " e.user_id STRING,"+
73 " IF(e.user_id=s.user_id"+
74 " OR (CAST(s.subscription_end_dt as BIGINT) < CAST(e.start_ts as BIGINT)), 'unsubscribed', 'subscribed') AS user_type,"+
75 " e.song_id STRING,"+
76 " e.artist_id STRING,"+
77 " (cast(e.end_ts as BIGINT)-cast(e.start_ts as BIGINT))/60 AS total_duration_in_minutes"+
78 " FROM project.enriched_data e"+
79 " LEFT OUTER JOIN project.subscribed_users s"+
80 " ON e.user_id=s.user_id"+
81 " WHERE e.status='pass'"+
82 " AND (batchid=$batchId)")
83

```

```

85 //<<<<<<----- PROBLEM 3 - Creation of table and Insertion of data ----->>>>>>>>
86 //Determine top 10 connected artists.
87 //Connected artists are those whose songs are most listened by the unique users who follow them.
88
89 val create_hive_table_top_10_connected_artists = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.connected_artists"+
90 " ("+
91 " artist_id STRING,"+
92 " total_distinct_songs INT,"+
93 " unique_followers INT"+
94 ")"+
95 " PARTITIONED BY (batchid INT)"+
96 " ROW FORMAT DELIMITED"+
97 " FIELDS TERMINATED BY ','"+
98 " STORED AS TEXTFILE")
99
100
101 val insert_into_top_10_connected_artists = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.connected_artists"+
102 " PARTITION (batchid=$batchId)"+
103 " SELECT"+
104 " artist_id,"+
105 " COUNT(DISTINCT song_id) AS total_distinct_songs,"+
106 " COUNT(DISTINCT user_id) AS unique_followers"+
107 " FROM project.enriched_data"+
108 " WHERE status='pass'"+
109 " AND (batchid=$batchId)"+
110 " GROUP BY artist_id"+
111 " ORDER BY unique_followers desc,total_distinct_songs desc"+
112 " LIMIT 10")
113
114

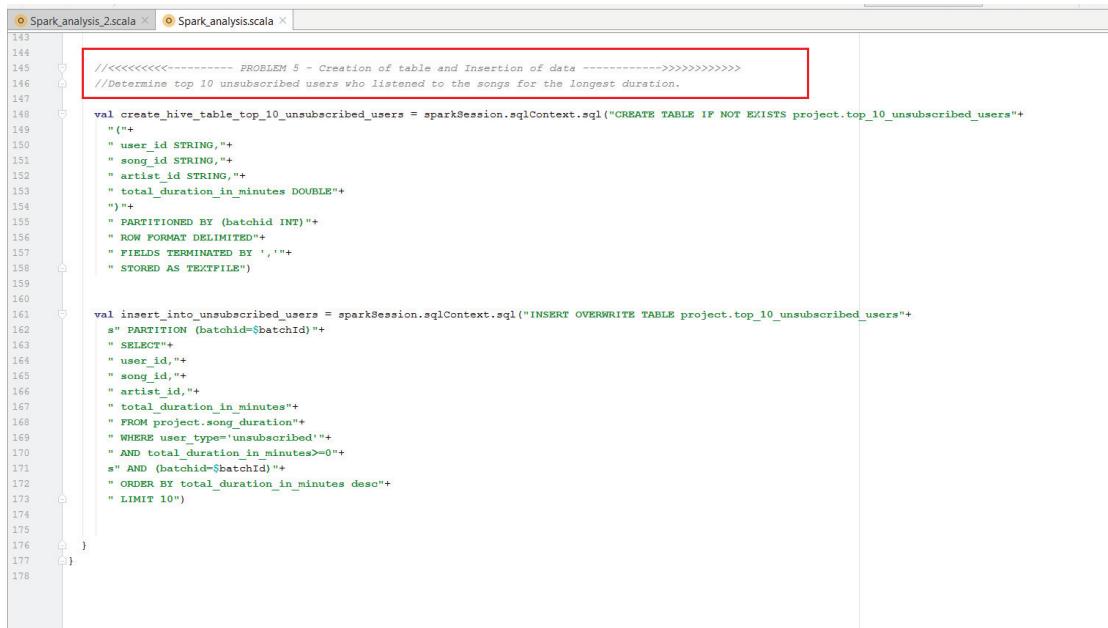
```

```

115 //<<<<<<----- PROBLEM 4 - Creation of table and Insertion of data ----->>>>>>>>
116 //Determine top 10 songs who have generated the maximum revenue.
117 //NOTE: Royalty applies to a song only if it was liked or was completed successfully or both.
118
119 val create_hive_table_top_10_songs_maxrevenue = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.top_10_songs_maxrevenue"+
120 " ("+
121 " song_id STRING,"+
122 " artist_id STRING,"+
123 " total_duration_in_minutes DOUBLE"+
124 " )"+
125 " PARTITIONED BY (batchid INT)"+
126 " ROW FORMAT DELIMITED"+
127 " FIELDS TERMINATED BY ','"+
128 " STORED AS TEXTFILE")
129
130
131 val insert_into_top_10_songs_maxrevenue = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.top_10_songs_maxrevenue"+
132 " PARTITION (batchid=$batchId)"+
133 " SELECT"+
134 " song_id,"+
135 " artist_id,"+
136 " (cast(end_ts as BIGINT)-cast(start_ts as BIGINT))/60 AS total_duration_in_minutes"+
137 " FROM project.enriched_data"+
138 " WHERE status='pass' +
139 " AND (batchid=$batchId)"+
140 " AND (like=1 OR song_end_type=0 OR (like=1 and song_end_type=0))"+
141 " ORDER BY total_duration_in_minutes desc"+
142 " LIMIT 10")
143

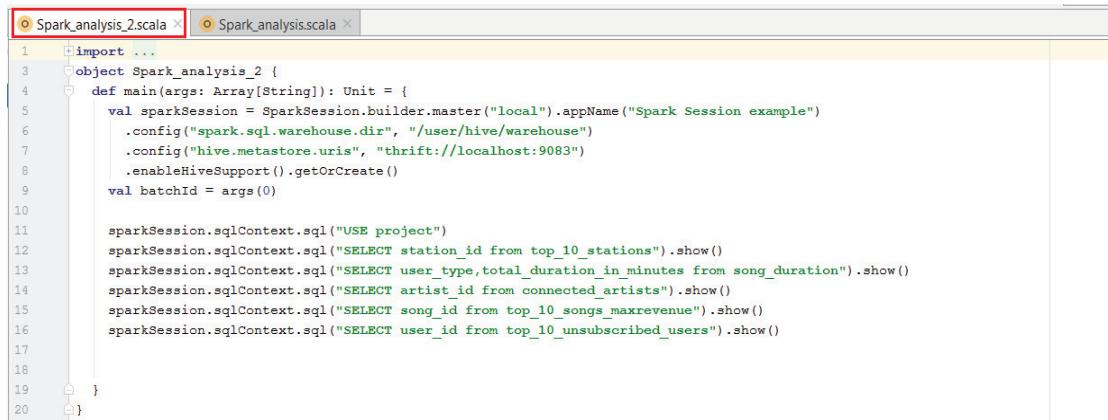
```

# Music Data Analysis



```
143 //<<<<<<----- PROBLEM 5 - Creation of table and Insertion of data ----->>>>>>>
144 //Determine top 10 unsubscribed users who listened to the songs for the longest duration.
145
146
147
148 val create_hive_table_top_10_unsubscribed_users = sparkSession.sqlContext.sql("CREATE TABLE IF NOT EXISTS project.top_10_unsubscribed_users"+
149
150 " ("+
151 " user_id STRING,"+
152 " song_id STRING,"+
153 " artist_id STRING,"+
154 " total_duration_in_minutes DOUBLE"+
155 ")"+
156 " PARTITIONED BY (batchid INT)"+
157 " ROW FORMAT DELIMITED"+
158 " FIELDS TERMINATED BY ','"+
159 " STORED AS TEXTFILE")
160
161 val insert_into_unsubscribed_users = sparkSession.sqlContext.sql("INSERT OVERWRITE TABLE project.top_10_unsubscribed_users"+
162 "s" PARTITION (batchid=$batchId)"+
163 " SELECT"+
164 " user_id,"+
165 " song_id,"+
166 " artist_id,"+
167 " total_duration_in_minutes"+
168 " FROM project.song_duration"+
169 " WHERE user_type='unsubscribed'"+
170 " AND total_duration_in_minutes>=0"+
171 "s" AND (batchid=$batchId)"+
172 " ORDER BY total_duration_in_minutes desc"+
173 " LIMIT 10")
174
175
176
177
178 }
```

This second Scala file is for displaying results of each query in the terminal.



```
1 import ...
2 object Spark_analysis_2 {
3   def main(args: Array[String]): Unit = {
4     val sparkSession = SparkSession.builder("local").appName("Spark Session example")
5       .config("spark.sql.warehouse.dir", "/user/hive/warehouse")
6       .config("hive.metastore.uris", "thrift://localhost:9083")
7       .enableHiveSupport().getOrCreate()
8     val batchId = args(0)
9
10    sparkSession.sqlContext.sql("USE project")
11    sparkSession.sqlContext.sql("SELECT station_id from top_10_stations").show()
12    sparkSession.sqlContext.sql("SELECT user_type, total_duration_in_minutes from song_duration").show()
13    sparkSession.sqlContext.sql("SELECT artist_id from connected_artists").show()
14    sparkSession.sqlContext.sql("SELECT song_id from top_10_songs_maxrevenue").show()
15    sparkSession.sqlContext.sql("SELECT user_id from top_10_unsubscribed_users").show()
16
17
18
19
20  }
```

- We will create jar for the above source code using File -> project structure -> Artifacts -> create jar using Build Artifacts command.
- Then we can observe **out** folder getting created with artifacts inside it.
- Now we will export this jar to VM where we will run using spark submit command and get the analysis result.

When we run the Data Analysis First it will submit Spark\_analysis file and then it will run the Spark-analysis\_2 file which displays the result for each query.

# Music Data Analysis

## Spark analysis output:

```
File Edit View Search Terminal Help
[acadgild@localhost scripts]$ ./DataAnalysis.sh
      zip warning: name not matched: META-INF/*.RSA
deleting: META-INF/DUMMY.SF
deleting: META-INF/DUMMY.DSA
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/Downloads/spark-2.0.0-bin-hadoop2.6/jars/slf4j-log4j12-1.7.16.jar!/org/slf4
/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
17/09/10 16:35:51 INFO SparkContext: Running Spark version 2.0.0
17/09/10 16:35:52 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classe
where applicable
17/09/10 16:35:52 WARN Utils: Your hostname, localhost.localdomain resolves to a loopback address: 127.0.0.1; using 192.168.
.4 instead (on interface eth5)
17/09/10 16:35:52 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
17/09/10 16:35:52 INFO SecurityManager: Changing view acls to: acadgild
17/09/10 16:35:52 INFO SecurityManager: Changing modify acls to: acadgild
17/09/10 16:35:52 INFO SecurityManager: Changing view acls groups to:
17/09/10 16:35:52 INFO SecurityManager: Changing modify acls groups to:
17/09/10 16:35:52 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users  with view permis
ry: /tmp/23a6594d-4ee0-422c-8777-11487adea113_resources
2017-09-10 16:56:59,310 INFO  [main] session.SessionState: Created HDFS director
y: /tmp/hive/acadgild/23a6594d-4ee0-422c-8777-11487adea113
2017-09-10 16:56:59,332 INFO  [main] session.SessionState: Created local directo
ry: /tmp/acadgild/23a6594d-4ee0-422c-8777-11487adea113
2017-09-10 16:56:59,365 INFO  [main] session.SessionState: Created HDFS director
y: /tmp/hive/acadgild/23a6594d-4ee0-422c-8777-11487adea113/_tmp_space.db
2017-09-10 16:56:59,374 INFO  [main] client.HiveClientImpl: Warehouse location f
or Hive client (version 1.2.1) is /user/hive/warehouse
2017-09-10 16:57:07,173 INFO  [main] execution.SparkSqlParser: Parsing command:
USE project
2017-09-10 16:57:07,332 INFO  [main] execution.SparkSqlParser: Parsing command:
CREATE TABLE IF NOT EXISTS project.top_10_stations( station_id STRING, total_dis
tinct_songs_played INT, distinct_user_count INT) PARTITIONED BY (batchid INT) RO
W FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE
2017-09-10 16:57:08,926 INFO  [main] execution.SparkSqlParser: Parsing command:
INSERT OVERWRITE TABLE project.top_10_stations PARTITION (batchid= 1) SELECT sta
tion_id, COUNT(DISTINCT song_id) AS total_distinct_songs_played, COUNT(DISTINCT
user_id) AS distinct_user_count FROM project.enriched_data WHERE status='pass' A
ND (batchid= 1) AND like=1 GROUP BY station_id ORDER BY total_distinct_songs_pla
yed DESC LIMIT 10
2017-09-10 16:57:11,668 INFO  [main] parser.CatalystSqlParser: Parsing command:
int
2017-09-10 16:57:11,689 INFO  [main] parser.CatalystSqlParser: Parsing command:
string
2017-09-10 16:57:11,711 INFO  [main] parser.CatalystSqlParser: Parsing command:
string
2017-09-10 16:57:11,713 INFO  [main] parser.CatalystSqlParser: Parsing command:
string
```

## Music Data Analysis

```
2017-09-10 16:58:57,921 INFO  [dag-scheduler-event-loop] scheduler.DAGScheduler: ResultStage 11 (sql at DataAnalysisMain_1.scala:131) finished in 1.228 s
2017-09-10 16:58:57,922 INFO  [main] scheduler.DAGScheduler: Job 4 finished: sql at DataAnalysisMain_1.scala:131, took 2.648835 s
2017-09-10 16:58:59,255 INFO  [main] common.FileUtils: deleting  hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_songs_maxrevenue/batchid=1/part-00000
2017-09-10 16:58:59,261 INFO  [main] fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
2017-09-10 16:58:59,337 INFO  [main] metadata.Hive: Replacing src:hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_songs_maxrevenue/.hive-staging_hive_2017-09-10_16-58-54_002 5741185161756501718-1/-ext-10000/part-00000, dest: hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_songs_maxrevenue/batchid=1/part-00000, Status:true
2017-09-10 16:58:59,833 INFO  [main] execution.SparkSqlParser: Parsing command: CREATE TABLE IF NOT EXISTS project.top_10_unsubscribed_users( user_id STRING, song_id STRING, artist_id STRING, total_duration_in_minutes DOUBLE) PARTITIONED BY (batchid INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE
2017-09-10 16:59:00,057 INFO  [main] execution.SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE project.top_10_unsubscribed_users PARTITION (batchid=1) SELECT user_id, song_id, artist_id, total_duration_in_minutes FROM project.song_duration WHERE user_type='unsubscribed' AND total_duration_in_minutes>=0 AND (batchid= 1) ORDER BY total_duration_in_minutes desc LIMIT 10
2017-09-10 16:59:00,490 INFO  [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 16:59:00,496 INFO  [main] parser.CatalystSqlParser: Parsing command: string
2017-09-10 16:59:00,496 INFO  [main] parser.CatalystSqlParser: Parsing command: string
```

```
2017-09-10 16:59:03,391 INFO  [dag-scheduler-event-loop] scheduler.DAGScheduler: Submitting 1 missing tasks from ResultStage 13 (MapPartitionsRDD[79] at sql at DataAnalysisMain_1.scala:161)
2017-09-10 16:59:03,392 INFO  [dag-scheduler-event-loop] scheduler.TaskSchedulerImpl: Adding task set 13.0 with 1 tasks
2017-09-10 16:59:03,394 INFO  [dispatcher-event-loop-0] scheduler.TaskSetManager: Starting task 0.0 in stage 13.0 (TID 809, localhost, partition 0, PROCESS_LOCAL, 5330 bytes)
2017-09-10 16:59:03,394 INFO  [Executor task launch worker-3] executor.Executor: Running task 0.0 in stage 13.0 (TID 809)
2017-09-10 16:59:03,483 INFO  [Executor task launch worker-3] storage.ShuffleBlockFetcherIterator: Getting 0 non-empty blocks out of 1 blocks
2017-09-10 16:59:03,483 INFO  [Executor task launch worker-3] storage.ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
2017-09-10 16:59:03,633 INFO  [Executor task launch worker-3] output.FileOutputCommitter: Saved output of task 'attempt_201709101659_0013_m_000000_0' to hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_unsubscribed_users/.hive-staging_hive_2017-09-10_16-59-01_939_5714149478602249311-1/-ext-10000/_temporary/0/task_201709101659_0013_m_000000
2017-09-10 16:59:03,633 INFO  [Executor task launch worker-3] mapred.SparkHadoopMapRedUtil: attempt_201709101659_0013_m_000000_0: Committed
2017-09-10 16:59:03,654 INFO  [Executor task launch worker-3] executor.Executor: Finished task 0.0 in stage 13.0 (TID 809). 1676 bytes result sent to driver
2017-09-10 16:59:03,655 INFO  [task-result-getter-1] scheduler.TaskSetManager: Finished task 0.0 in stage 13.0 (TID 809) in 262 ms on localhost (1/1)
2017-09-10 16:59:03,655 INFO  [task-result-getter-1] scheduler.TaskSchedulerImpl: Removed TaskSet 13.0, whose tasks have all completed, from pool
2017-09-10 16:59:03,655 INFO  [dag-scheduler-event-loop] scheduler.DAGScheduler: ResultStage 13 (sql at DataAnalysisMain_1.scala:161) finished in 0.254 s
2017-09-10 16:59:03,657 INFO  [main] scheduler.DAGScheduler: Job 5 finished: sql at DataAnalysisMain_1.scala:161, took 0.747655 s
```

## Music Data Analysis

```
2017-09-10 18:07:03,355 INFO [main] session.SessionState: Created HDFS directory: /tmp/hive/acadgild/d62b1207-5c3f-4105-b4cd-7b887ebda0f7
2017-09-10 18:07:03,395 INFO [main] session.SessionState: Created local directory: /tmp/acadgild/d62b1207-5c3f-4105-b4cd-7b887ebda0f7
2017-09-10 18:07:03,421 INFO [main] session.SessionState: Created HDFS directory: /tmp/hive/acadgild/d62b1207-5c3f-4105-b4cd-7b887ebda0f7/_tmp_space.db
2017-09-10 18:07:03,436 INFO [main] client.HiveClientImpl: Warehouse location for Hive client (version 1.2.1) is /user/hive/warehouse
2017-09-10 18:07:12,732 INFO [main] execution.SparkSqlParser: Parsing command: USE project
2017-09-10 18:07:13,108 INFO [main] execution.SparkSqlParser: Parsing command: select station_id from top_10_stations
2017-09-10 18:07:19,846 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:19,910 INFO [main] parser.CatalystSqlParser: Parsing command: string
2017-09-10 18:07:19,920 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:19,921 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:20,328 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:20,329 INFO [main] parser.CatalystSqlParser: Parsing command: string
2017-09-10 18:07:20,330 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:20,330 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:07:24,383 INFO [main] memory.MemoryStore: Block broadcast_0 stored as values in memory (estimated size 148.4 KB, free 413.8 MB)
2017-09-10 18:07:24,987 INFO [main] memory.MemoryStore: Block broadcast_0_piece
```

### Spark analysis\_2 output:

Query-1: Determine top 10 station\_id(s) where maximum number of songs were played, which were liked by unique users.

```
2017-09-10 18:21:10,509 INFO [main] scheduler.DAGScheduler: Job 0 finished: show at DataAnalysisReadFromHive.scala:22, took 15.957828 s
2017-09-10 18:21:10,994 INFO [main] codegen.CodeGenerator: Code generated in 19.9.417505 ms
+-----+
|station_id|
+-----+
|  ST402|
|  ST411|
|  ST405|
|  ST410|
+-----+
2017-09-10 18:21:11,227 INFO [main] execution.SparkSqlParser: Parsing command: select user_type,sum(total_duration_in_minutes) as total_song_duration from song_duration where total_duration_in_minutes>=0 group by user_type
2017-09-10 18:21:11,715 INFO [main] parser.CatalystSqlParser: Parsing command: int
```

Query-2: Determine total duration of songs played by each type of user, where type of user can be 'subscribed' or 'unsubscribed'. An unsubscribed user is the one whose record is either not present in Subscribed\_users lookup table or has subscription\_end\_date earlier than the timestamp of the song played by him.

```
finished task 186.0 in stage 4.0 (TID 201) in 470 ms on localhost (199/199)
2017-09-10 18:21:37,544 INFO [task-result-getter-1] scheduler.TaskSchedulerImpl: Removed TaskSet 4.0, whose tasks have all completed, from pool
2017-09-10 18:21:37,546 INFO [main] scheduler.DAGScheduler: Job 2 finished: show at DataAnalysisReadFromHive.scala:31, took 14.650864 s
2017-09-10 18:21:37,705 INFO [main] codegen.CodeGenerator: Code generated in 86.765409 ms
+-----+-----+
| user_type|total_song_duration|
+-----+-----+
|subscribed| 1904665.6500000001|
+-----+-----+
2017-09-10 18:21:37,721 INFO [main] execution.SparkSqlParser: Parsing command: select artist_id from connected_artists
2017-09-10 18:21:37,902 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:21:37,909 INFO [main] parser.CatalystSqlParser: Parsing command: string
```

## Music Data Analysis

Query-3: Determine top 10 connected artists. Connected artists are those whose songs are most listened by the unique users who follow them

```
2017-09-10 18:21:39,401 INFO [dag-scheduler-event-loop] scheduler.DAGScheduler: ResultStage 5 (show at DataAnalysisReadFromHive.scala:37) finished in 0.224 s
2017-09-10 18:21:39,404 INFO [main] scheduler.DAGScheduler: Job 3 finished: show at DataAnalysisReadFromHive.scala:37, took 0.297580 s
2017-09-10 18:21:39,407 INFO [task-result-getter-2] scheduler.TaskSchedulerImpl: Removed TaskSet 5.0, whose tasks have all completed, from pool
+-----+
|artist_id|
+-----+
|  A300|
|  A303|
|  A304|
|  A305|
|  A302|
+-----+
2017-09-10 18:21:39,421 INFO [main] execution.SparkSqlParser: Parsing command: select song_id from top_10_songs_maxrevenue
2017-09-10 18:21:40,225 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:21:40,226 INFO [main] parser.CatalystSqlParser: Parsing command: string
```

Query-4: Determine top 10 songs who have generated the maximum revenue. Royalty applies to a song only if it was liked or was completed successfully or both

```
finished task 0.0 in stage 6.0 (TID 203) in 155 ms on localhost (1/1)
2017-09-10 18:21:41,756 INFO [task-result-getter-3] scheduler.TaskSchedulerImpl: Removed TaskSet 6.0, whose tasks have all completed, from pool
+-----+
|song_id|
+-----+
|  S209|
|  S202|
|  S205|
|  S200|
|  S203|
|  S203|
|  S206|
|  S202|
|  S206|
|  S202|
+-----+
2017-09-10 18:21:41,774 INFO [main] execution.SparkSqlParser: Parsing command: select user_id from top_10_unsubscribed_users
2017-09-10 18:21:41,880 INFO [main] parser.CatalystSqlParser: Parsing command: int
2017-09-10 18:21:41,882 INFO [main] parser.CatalystSqlParser: Parsing command: string
```

Query-5: Determine top 10 unsubscribed users who listened to the songs for the longest duration.

```
2017-09-10 18:21:43,240 INFO [task-result-getter-0] scheduler.TaskSetManager: F
finished task 0.0 in stage 7.0 (TID 204) in 140 ms on localhost (1/1)
2017-09-10 18:21:43,241 INFO [task-result-getter-0] scheduler.TaskSchedulerImpl: Removed TaskSet 7.0, whose tasks have all completed, from pool
+-----+
|user_id|
+-----+
2017-09-10 18:21:43,356 INFO [Thread-2] spark.SparkContext: Invoking stop() fro
m shutdown hook
2017-09-10 18:21:43,576 INFO [Thread-2] server.ServerConnector: Stopped ServerC
onnection@ff8277e{HTTP/1.1}{0.0.0.0:4040}
```

For this query we got no result because there is no unsubscribed user who has got positive total\_duration\_minutes\_in\_minutes. All the unsubscribed users got negative total\_duration means their end timestamp is less than the start time stamp of the song which is not correct.

## Music Data Analysis

Now we can see the output of Spark analysis in hive:

```
hive> show tables;
OK
enriched_data
formatted_input
song_artist_map
station_geo_map
subscribed_users
users_artists

Time taken: 0.04 seconds, Fetched: 6 row(s)
hive> show tables;
OK
connected_artists
enriched_data
formatted_input
song_artist_map
song_duration
station_geo_map
subscribed_users
top_10_songs_maxrevenue
top_10_stations
top_10_unsubscribed_users
users_artists

Time taken: 0.047 seconds, Fetched: 11 row(s)
hive> █
```

```
hive> select * from song_duration;
OK
U115    subscribed    S200    A300    -480116.766666666666    1
U108    subscribed    S200    A302    87193.7833333334    1
U120    subscribed    S201    A300    43405.55    1
U107    unsubscribed  S202    A304    -87193.7833333334    1
U101    subscribed    S202    A300    166666.6666666666    1
U110    unsubscribed  S202    A303    -436711.2166666667    1
U118    subscribed    S202    A300    0.0    1
U104    subscribed    S202    A303    -166666.6666666666    1
U102    subscribed    S203    A305    480116.7666666666    1
U113    subscribed    S203    A304    -43788.2333333333    1
U113    subscribed    S203    A303    436711.2166666667    1
U105    subscribed    S203    A303    -1666.666666666667    1
U113    subscribed    S203    A303    0.0    1
U103    unsubscribed  S204    A300    -480116.766666666666    1
U108    subscribed    S205    A304    166666.6666666666    1
U106    subscribed    S206    A300    -43788.2333333333    1
U120    subscribed    S206    A303    -333333.3333333333    1
U105    unsubscribed  S207    A300    -333333.3333333333    1
U116    subscribed    S208    A303    -166666.6666666666    1
U114    subscribed    S209    A303    523905.0    1
Time taken: 0.048 seconds, Fetched: 20 row(s)
hive> select * from top_10_stations;
OK
ST402    2    2    1
ST411    2    2    1
ST405    1    1    1
ST410    1    1    1
Time taken: 0.061 seconds, Fetched: 4 row(s)
hive> █
```

## Music Data Analysis

```

hive> select * from top_10_songs_maxrevenue;
OK
S209    A303    523905.0      1
S202    A300    166666.6666666666  1
S205    A304    166666.6666666666  1
S200    A302    87193.7833333334  1
S203    A303    0.0      1
S203    A304    -43788.2333333333  1
S206    A300    -43788.2333333333  1
S202    A304    -87193.7833333334  1
S206    A303    -333333.333333333  1
S202    A303    -436711.2166666667  1
Time taken: 0.063 seconds, Fetched: 10 row(s)
hive> select * from connected_artist;
FAILED: SemanticException [Error 10001]: Line 1:14 Table not found 'connected_artist'
hive> select * from connected_artists;
OK
A300    6      7      1
A303    5      7      1
A304    3      3      1
A305    1      1      1
A302    1      1      1
Time taken: 0.062 seconds, Fetched: 5 row(s)
hive> select * from top_10_unsubscribed_users;
OK
Time taken: 0.057 seconds
hive> ■

```

```

hive> select user_type, SUM(total_duration_in_minutes) from song_duration where total_duration_in_minutes >=0 group by user_type;
Query ID = acadgild_20170911202020_14573a7d-939c-410f-aff0-0de67d02cbc2
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1505113626668_0001, Tracking URL = http://localhost:8088/proxy/application_1505113626668_0001/
Kill Command = /home/acadgild/hadoop-2.6.0/bin/hadoop job -kill job_1505113626668_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2017-09-11 20:20:59,578 Stage-1 map = 0%, reduce = 0%
2017-09-11 20:21:07,222 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.95 sec
2017-09-11 20:21:15,787 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.85 sec
MapReduce Total cumulative CPU time: 4 seconds 850 msec
Ended Job = job_1505113626668_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.85 sec HDFS Read: 1102 HDFS Write: 30 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 850 msec
OK
subscribed    1904665.6500000001
Time taken: 30.85 seconds, Fetched: 1 row(s)
hive> ■

```

Now we have seen all the spark queries creating the tables for each query. So Data Analysis using Spark is executed successfully.

### g) Data Export to MYSQL:

In this stage data will be exported from hive warehouse directory to MYSQL database using data\_export.sh command mentioned in the DataAnalysis.sh script.

```

#!/bin/bash
batchid=`cat /home/acadgild/project/logs/current-batch.txt`
LOGFILE=/home/acadgild/project/logs/log_batch_$batchid
echo "Running script for data analysis using spark..." >> $LOGFILE
chmod 775 /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar
zip -d /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar META-INF/*.DSA META-INF/*.RSA META-INF/*.SF
spark-submit \
--class Spark_analysis \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid
spark-submit \
--class Spark_analysis_2 \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid
echo "Exporting data to MYSQL using sqoop export" >> $LOGFILE
sh /home/acadgild/project/scripts/data_export.sh
echo "Incrementing batchid..." >> $LOGFILE
batchid=`expr $batchid + 1`
echo -n $batchid > /home/acadgild/project/logs/current-batch.txt

```

## Music Data Analysis

```
#!/bin/bash
batchid=`cat /home/acadgild/project/logs/current-batch.txt`
LOGFILE=/home/acadgild/project/logs/log_batch_$batchid
echo "Creating mysql tables if not present..." >> $LOGFILE
mysql -u root </home/acadgild/project/scripts/create_schema.sql
echo "Running sqoop job for data export..." >> $LOGFILE

sqoop export \
--connect jdbc:mysql://localhost/project \
--username 'root' \
--P \
--table 'top_10_stations' \
--export-dir '/user/hive/warehouse/project.db/top_10_stations/batchid=$batchid/part-00000' \
--input-fields-terminated-by ',' \
-m 1

sqoop export \
--connect jdbc:mysql://localhost/project \
--username 'root' \
--P \
--table 'song_duration' \
--export-dir '/user/hive/warehouse/project.db/song_duration/batchid=$batchid/part-00000' \
--input-fields-terminated-by ',' \
-m 1

sqoop export \
--connect jdbc:mysql://localhost/project \
--username 'root' \
--P \
--table 'connected_artists' \
--export-dir '/user/hive/warehouse/project.db/connected_artists/batchid=$batchid/part-00000' \
--input-fields-terminated-by ',' \
-m 1
```

```
sqoop export \
--connect jdbc:mysql://localhost/project \
--username 'root' \
--P \
--table 'top_10_songs_maxrevenue' \
--export-dir '/user/hive/warehouse/project.db/top_10_songs_maxrevenue/batchid=$batchid/part-00000' \
--input-fields-terminated-by ',' \
-m 1

sqoop export \
--connect jdbc:mysql://localhost/project \
--username 'root' \
--P \
--table 'top_10_unsubscribed_users' \
--export-dir '/user/hive/warehouse/project.db/top_10_unsubscribed_users/batchid=$batchid/part-00000' \
--input-fields-terminated-by ',' \
-m 1
```

```
Warning: /usr/local/sqoop/../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/../zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-09-10 23:00:16,788 INFO  [main] sqoop.Sqoop: Running Sqoop version: 1.4.5
Enter password:
2017-09-10 23:00:18,522 INFO  [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-09-10 23:00:18,523 INFO  [main] tool.CodeGenTool: Beginning code generation
2017-09-10 23:00:18,751 INFO  [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `song_duration` AS t LIMIT 1
2017-09-10 23:00:18,770 INFO  [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `song_duration` AS t LIMIT 1
2017-09-10 23:00:18,776 INFO  [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0
Note: /tmp/sqoop-acadgild/compile/d4aa31e4f5d01298a20c5991eb23795c/song_duration.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2017-09-10 23:00:19,955 INFO  [main] orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/d4aa31e4f5d01298a20c5991eb23795c/song_duration.jar
2017-09-10 23:00:19,963 INFO  [main] mapreduce.ExportJobBase: Beginning export of song_duration
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-09-10 23:00:20,154 WARN  [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using built-in java classes where applicable
2017-09-10 23:00:20,159 INFO  [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
2017-09-10 23:00:20,764 INFO  [main] Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
2017-09-10 23:00:20,767 INFO  [main] Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
2017-09-10 23:00:20,768 INFO  [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.map
```

## Music Data Analysis

```
File Edit View Search Terminal Help
2017-09-10 23:00:21,836 INFO [main] Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapreduce.inputformat.class is deprecated. Instead, use mapreduce.job.inputformat.class
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.mapoutput.value.class is deprecated. Instead, use mapreduce.map.output.value.class
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.mapoutput.key.class is deprecated. Instead, use mapreduce.map.output.key.class
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: mapred.job.classpath.files is deprecated. Instead, use mapreduce.job.classpath.files
2017-09-10 23:00:21,837 INFO [main] Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name
2017-09-10 23:00:21,838 INFO [main] Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2017-09-10 23:00:21,838 INFO [main] Configuration.deprecation: mapred.cache.files.filesizes is deprecated. Instead, use mapreduce.job.cache.files.filesizes
2017-09-10 23:00:21,942 INFO [main] mapreduce.JobSubmitter: Submitting tokens for job: job_1505059429035_0009
2017-09-10 23:00:22,143 INFO [main] impl.YarnClientImpl: Submitted application application_1505059429035_0009 to ResourceManager at /0.0.0.0:8032
2017-09-10 23:00:22,184 INFO [main] mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1505059429035_0009/
2017-09-10 23:00:22,184 INFO [main] mapreduce.Job: Running job: job_1505059429035_0009
2017-09-10 23:00:28,333 INFO [main] mapreduce.Job: Job job_1505059429035_0009 running in uber mode : false
2017-09-10 23:00:28,336 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-09-10 23:00:34,425 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-09-10 23:00:34,435 INFO [main] mapreduce.Job: Job job_1505059429035_0009 completed successfully
```

Now we can see the data exported successfully into the MYSQL Database for all the 5 queries.

MYSQL:

```
File Edit View Search Terminal Help
mysql> use project;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+---------------------+
| Tables_in_project |
+-----+
| connected_artists |
| song_duration      |
| top_10_royalty_songs |
| top_10_stations    |
| top_10_unsubscribed_users |
+-----+
5 rows in set (0.00 sec)

mysql> select * from connected_artists;
+-----+-----+-----+
| artist_id | total_distinct_songs | user_count |
+-----+-----+-----+
| A300     | 6 | 7 |
| A303     | 5 | 7 |
| A304     | 3 | 3 |
| A305     | 1 | 1 |
| A302     | 1 | 1 |
+-----+-----+-----+
5 rows in set (0.05 sec)
```

## Music Data Analysis

```
File Edit View Search Terminal Help
mysql> select * from top_10_royalty_songs;
+-----+-----+-----+
| song_id | artist_id | duration |
+-----+-----+-----+
| S209 | A303 | 523905 |
| S202 | A300 | 166666.66666667 |
| S205 | A304 | 166666.66666667 |
| S200 | A302 | 87193.7833333333 |
| S203 | A303 | 0 |
| S203 | A304 | -43788.2333333333 |
| S206 | A300 | -43788.2333333333 |
| S202 | A304 | -87193.7833333333 |
| S206 | A303 | -333333.333333333 |
| S202 | A303 | -436711.216666667 |
+-----+-----+-----+
10 rows in set (0.01 sec)

mysql> select * from top_10_unsubscribed_users;
Empty set (0.00 sec)

mysql> select * from top_10_stations;
+-----+-----+-----+
| station_id | total_distinct_songs_played | distinct_user_count |
+-----+-----+-----+
| ST402 | 2 | 2 |
| ST411 | 2 | 2 |
| ST405 | 1 | 1 |
| ST410 | 1 | 1 |
+-----+-----+-----+
4 rows in set (0.02 sec)
```

```
File Edit View Search Terminal Help
mysql> select * from top_10_stations;
+-----+-----+-----+
| station_id | total_distinct_songs_played | distinct_user_count |
+-----+-----+-----+
| ST402 | 2 | 2 |
| ST411 | 2 | 2 |
| ST405 | 1 | 1 |
| ST410 | 1 | 1 |
+-----+-----+-----+
4 rows in set (0.02 sec)

mysql> select * from song_duration;
+-----+-----+-----+-----+
| user_id | user_type | song_id | artist_id | total_duration |
+-----+-----+-----+-----+
| U115 | subscribed | S200 | A300 | -480116.766666667 |
| U108 | subscribed | S200 | A302 | 87193.7833333333 |
| U120 | subscribed | S201 | A300 | 43405.55 |
| U107 | unsubscribed | S202 | A304 | -87193.7833333333 |
| U101 | subscribed | S202 | A300 | 166666.666666667 |
| U110 | unsubscribed | S202 | A303 | -436711.216666667 |
| U118 | subscribed | S202 | A300 | 0 |
| U104 | subscribed | S202 | A303 | -166666.666666667 |
| U102 | subscribed | S203 | A305 | 480116.766666667 |
| U113 | subscribed | S203 | A304 | -43788.2333333333 |
| U113 | subscribed | S203 | A303 | 436711.216666667 |
| U105 | subscribed | S203 | A303 | -1666.66666666667 |
| U113 | subscribed | S203 | A303 | 0 |
| U103 | unsubscribed | S204 | A300 | -480116.766666667 |
| U108 | subscribed | S205 | A304 | 166666.666666667 |
| U106 | subscribed | S206 | A300 | -43788.2333333333 |
| U120 | subscribed | S206 | A303 | -333333.333333333 |
| U105 | unsubscribed | S207 | A300 | -333333.333333333 |
| U116 | subscribed | S208 | A303 | -166666.666666667 |
| U114 | subscribed | S209 | A303 | 523905 |
+-----+-----+-----+-----+
20 rows in set (0.02 sec)

mysql> 
```

```
acadgild@localhost:~ | start-daemon.sh (~/D... | acadgild@localhost:~
```

```
mysql> select user_type,sum(total_duration) from song_duration where total_duration >=0 group by user_type;
+-----+-----+
| user_type | sum(total_duration) |
+-----+-----+
| subscribed | 1904665.65 |
+-----+-----+
1 row in set (0.00 sec)
```

```
mysql> 
```

## Music Data Analysis

```
start-daemon.sh data_export.sh log_batch_1 DataAnalysis.sh wrapper.sh log_batch_2
```

```
#!/bin/bash

batchid=`cat /home/acadgild/project/logs/current-batch.txt`

LOGFILE=/home/acadgild/project/logs/log_batch_$batchid

echo "Running script for data analysis using spark..." >> $LOGFILE
chmod 775 /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar

zip -d /home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar META-INF/*.DSA META-INF/*.SF

spark-submit \
--class Spark_analysis \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid

spark-submit \
--class Spark_analysis_2 \
--master local[2] \
--driver-class-path /usr/local/hive/lib/hive-hbase-handler-0.14.0.jar:/usr/local/hbase/lib/* \
/home/acadgild/project/Spark_Practice-1/out/artifacts/Spark_practice_1_jar/Spark_practice-1.jar $batchid

echo "Exporting data to MySQL using sqoop export..." >> $LOGFILE
sh /home/acadgild/project/scripts/data_export.sh

echo "Incrementing batchid..." >> $LOGFILE
batchid= `expr $batchid + 1`
echo -n $batchid > /home/acadgild/project/logs/current-batch.txt
```

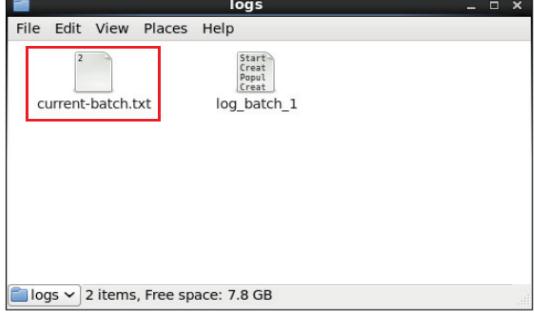
Now after exporting data into MySQL batchid will be incremented to 2 means one batch of data operations is successfully completed and new batch of data will be loaded for the analysis after every 3 hours.

```
[acadgild@localhost scripts]$ cat /home/acadgild/project/logs/current-batch.txt
2[acadgild@localhost scripts]$
```

We can check logs to track the behaviour of the operations we have done on the data and overcome failures in the pipeline and we can see the batchid incremented value in current-batch.txt.

```
start-daemon.sh data_export.sh log_batch_1 DataAnalysis.sh wrapper.sh
```

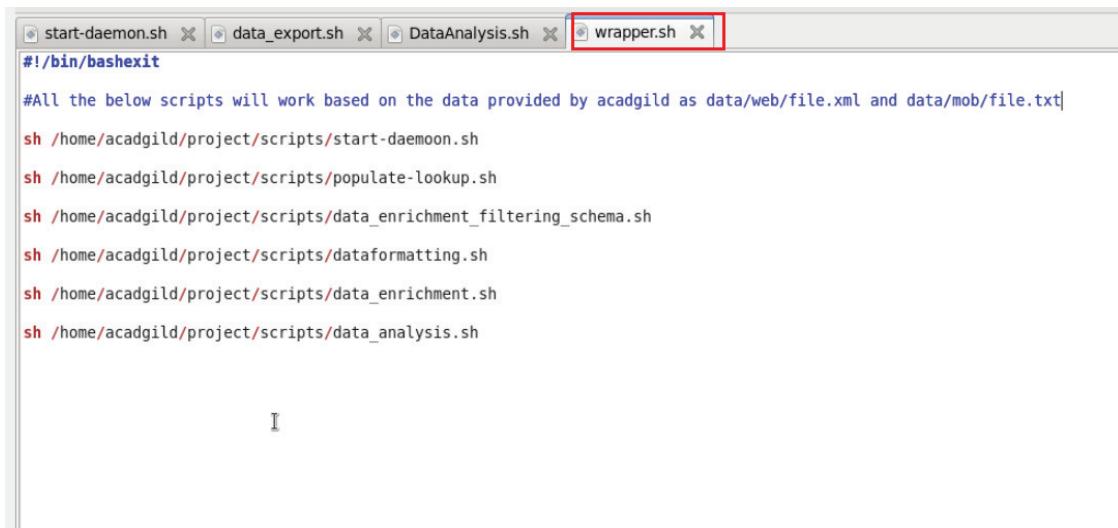
```
Starting daemons...
Creating Lookup Tables...
Populating Lookup Tables...
Creating hive tables on top hbase tables for data enrichment and filtering...
Placing data files from local to HDFS...
Running pig script for data formatting...
Running hive script for formatted data load...
Running hive script for data enrichment and filtering...
Copying valid and invalid records in local file system...
Deleting older valid and invalid records from local file system...
Running script for data analysis using spark...
Creating mysql tables if not present...
Exporting data to MySQL using sqoop export...
Running sqoop job for data export...
Incrementing batchid...
```



The screenshot shows a file browser window titled "logs". Inside the window, there are two files: "current-batch.txt" and "log\_batch\_1". The "current-batch.txt" file is highlighted with a red box. The window has a standard menu bar with File, Edit, View, Places, Help, and a toolbar with icons for Start, Create, Popul, and Creat.

## Music Data Analysis

- Wrapping all the scripts inside the single script file and scheduling this file to run at the periodic interval of every 3 hours.

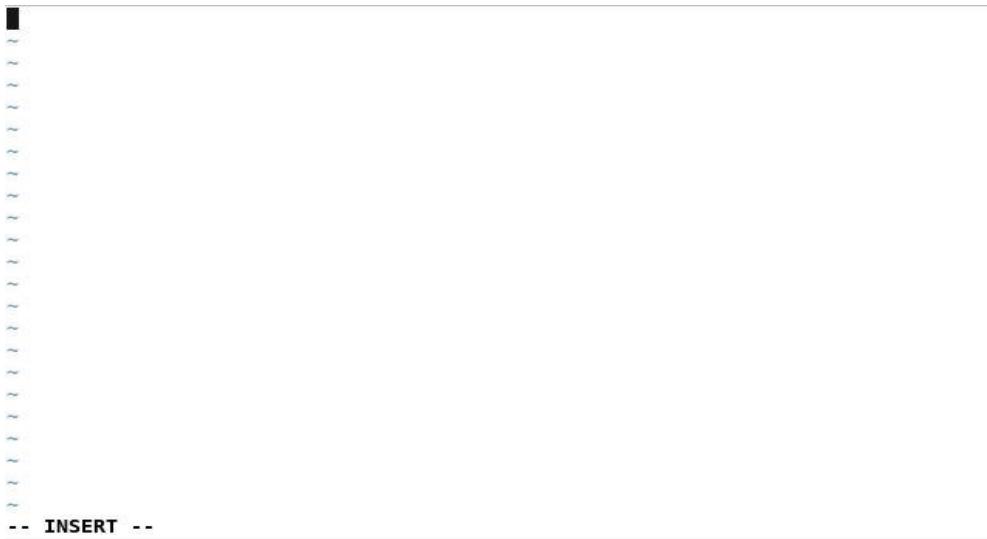


```
#!/bin/bashexit
>All the below scripts will work based on the data provided by acadgild as data/web/file.xml and data/mob/file.txt|
sh /home/acadgild/project/scripts/start-daemon.sh
sh /home/acadgild/project/scripts/populate-lookup.sh
sh /home/acadgild/project/scripts/data_enrichment_filtering_schema.sh
sh /home/acadgild/project/scripts/dataformatting.sh
sh /home/acadgild/project/scripts/data_enrichment.sh
sh /home/acadgild/project/scripts/data_analysis.sh
```

### h) Job Scheduling:

- Creating Crontab to schedule the wrapper.sh script to run for every 3 hour interval.

```
[acadgild@localhost ~]$ cd project
[acadgild@localhost project]$ crontab -e
```



```
-- INSERT --
```



## Music Data Analysis

### 7. Conclusion:

So we performed all the data operations as per the sequence mentioned in the wrapper.sh file and obtained results successfully for the one of the leading music company.